

Wanat_Assignment2_final

July 7, 2019

```
In [1]: # Jump-Start for the Bank Marketing Study
# as described in Marketing Data Science: Modeling Techniques
# for Predictive Analytics with R and Python (Miller 2015)

# jump-start code revised by Thomas W. Milller (2018/10/07)

# Scikit Learn documentation for this assignment:
# http://scikit-learn.org/stable/auto_examples/classification/
#   plot_classifier_comparison.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.naive_bayes.BernoulliNB.html#sklearn.naive_bayes.BernoulliNB.score
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.linear_model.LogisticRegression.html
# http://scikit-learn.org/stable/modules/model_evaluation.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.model_selection.KFold.html

# prepare for Python version 3x features and functions
# comment out for Python 3.x execution
# from __future__ import division, print_function
# from future_builtins import ascii, filter, hex, map, oct, zip

# seed value for random number generators to obtain reproducible results
RANDOM_SEED = 1

# import base packages into the namespace for this program
import numpy as np
import pandas as pd
import pandas_profiling
import matplotlib
import matplotlib.pyplot as plt # static plotting
import seaborn as sns # pretty plotting, including heat map
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import classification_report
from sklearn.metrics import roc_curve
```

```

from sklearn.preprocessing import binarize
from sklearn.metrics import precision_recall_curve
from sklearn.naive_bayes import BernoulliNB
#from sklearn.cross_validation import cross_val_score
from sklearn.model_selection import cross_val_score
import math

```

0.1 Defined functions

```

In [2]: # correlation heat map setup for seaborn
def corr_chart(df_corr):
    corr=df_corr.corr()
    #screen top half to get a triangle
    top = np.zeros_like(corr, dtype=np.bool)
    top[np.triu_indices_from(top)] = True
    fig=plt.figure()
    fig, ax = plt.subplots(figsize=(12,12))
    sns.heatmap(corr, mask=top, cmap='coolwarm',
                center = 0, square=True,
                linewidths=.5, cbar_kws={'shrink':.5},
                annot = True, annot_kws={'size': 9}, fmt = '.3f')
    plt.xticks(rotation=45) # rotate variable labels on columns (x axis)
    plt.yticks(rotation=0) # use horizontal variable labels on rows (y axis)
    plt.title('Correlation Heat Map')
    plt.savefig('plot-corr-map.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)

np.set_printoptions(precision=3)

In [3]: # define a function to return model metrics for evaluation

def model_metrics(y_known, y_pred):
    y_test = y_known
    y_pred_class = y_pred
    confusion = metrics.confusion_matrix(y_test, y_pred_class)
    TP = confusion[1, 1]
    TN = confusion[0, 0]
    FP = confusion[0, 1]
    FN = confusion[1, 0]
    accuracy = metrics.accuracy_score(y_test, y_pred_class)
    class_error = 1 - metrics.accuracy_score(y_test, y_pred_class)
    sensitivity = metrics.recall_score(y_test, y_pred_class)
    specificity = TN / (TN + FP)
    false_positive_rate = FP / float(TN + FP)
    precision = TP / float(TP + FP)
    f1_score = metrics.f1_score(y_test, y_pred_class)

```

```

#      $F1 = 2 * precision * sensitivity / (precision + sensitivity)$ 
print('The model metrics are:',
      '\naccuracy:', accuracy,
      '\nclassification error:', class_error,
      '\nsensitivity:', sensitivity,
      '\nspecificity:', specificity,
      '\nfalse positive rate:', false_positive_rate,
      '\nprecision:', precision,
      '\nF1 score:', f1_score,
#       '\nF1 by hand:', F1,
      '\nconfusion matrix:')
return(confusion)

```

In [4]: # define a function that accepts a threshold and
prints sensitivity and specificity

```

def evaluate_threshold(threshold):
    print('Sensitivity:', tpr[thresholds > threshold][-1])
    print('Specificity:', 1 - fpr[thresholds > threshold][-1])

```

In [5]: # define a function that accepts the fpr and tpr values
from the roc_curve function and plot the ROC curve

```

def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'k--')
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC curve for response')

```

In [6]: # define a function that accepts the recall and precision values
from the precision_recall_curve function and plot the PR curve

```

def plot_pr_curve(recall, precision, label=None):
    plt.plot(recall, precision, linewidth=2, label=label)
#     plt.plot([0, 1], [0, 1], 'k--')
    plt.axis([0, 1, 0, 1])
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.title('PR curve for response')

```

In [7]: # define a function that accepts the scores from the cross validation
and print the scores, mean, and standard deviation

```

def display_scores(scores):
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard deviation:", scores.std())

```

0.2 Import data set and prepare for analysis

```
In [8]: # initial work with the smaller data set
        bank = pd.read_csv('bank.csv', sep = ';') # start with smaller data set
        # examine the shape of original input data
        print(bank.shape)
```

(4521, 17)

```
In [9]: #total number of NaN values in each column
        bank.isnull().sum()
```

```
Out[9]: age          0
        job          0
        marital      0
        education    0
        default      0
        balance      0
        housing      0
        loan         0
        contact      0
        day          0
        month        0
        duration     0
        campaign     0
        pdays       0
        previous     0
        poutcome     0
        response     0
        dtype: int64
```

```
In [10]: print(bank.response.head())
```

```
0    no
1    no
2    no
3    no
4    no
```

Name: response, dtype: object

```
In [11]: display(bank.head())
```

	age	job	marital	education	default	balance	housing	loan	\
0	30	unemployed	married	primary	no	1787	no	no	
1	33	services	married	secondary	no	4789	yes	yes	
2	35	management	single	tertiary	no	1350	yes	no	
3	30	management	married	tertiary	no	1476	yes	yes	

```
4 59 blue-collar married secondary no 0 yes no
```

	contact	day	month	duration	campaign	pdays	previous	poutcome	response
0	cellular	19	oct	79	1	-1	0	unknown	no
1	cellular	11	may	220	1	339	4	failure	no
2	cellular	16	apr	185	1	330	1	failure	no
3	unknown	3	jun	199	4	-1	0	unknown	no
4	unknown	5	may	226	1	-1	0	unknown	no

```
In [12]: bank.describe()
```

```
Out[12]:
```

	age	balance	day	duration	campaign	\
count	4521.000000	4521.000000	4521.000000	4521.000000	4521.000000	
mean	41.170095	1422.657819	15.915284	263.961292	2.793630	
std	10.576211	3009.638142	8.247667	259.856633	3.109807	
min	19.000000	-3313.000000	1.000000	4.000000	1.000000	
25%	33.000000	69.000000	9.000000	104.000000	1.000000	
50%	39.000000	444.000000	16.000000	185.000000	2.000000	
75%	49.000000	1480.000000	21.000000	329.000000	3.000000	
max	87.000000	71188.000000	31.000000	3025.000000	50.000000	

	pdays	previous
count	4521.000000	4521.000000
mean	39.766645	0.542579
std	100.121124	1.693562
min	-1.000000	0.000000
25%	-1.000000	0.000000
50%	-1.000000	0.000000
75%	-1.000000	0.000000
max	871.000000	25.000000

```
In [13]: bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4521 entries, 0 to 4520
Data columns (total 17 columns):
age          4521 non-null int64
job          4521 non-null object
marital      4521 non-null object
education    4521 non-null object
default      4521 non-null object
balance      4521 non-null int64
housing      4521 non-null object
loan         4521 non-null object
contact      4521 non-null object
day          4521 non-null int64
month        4521 non-null object
duration     4521 non-null int64
```

```

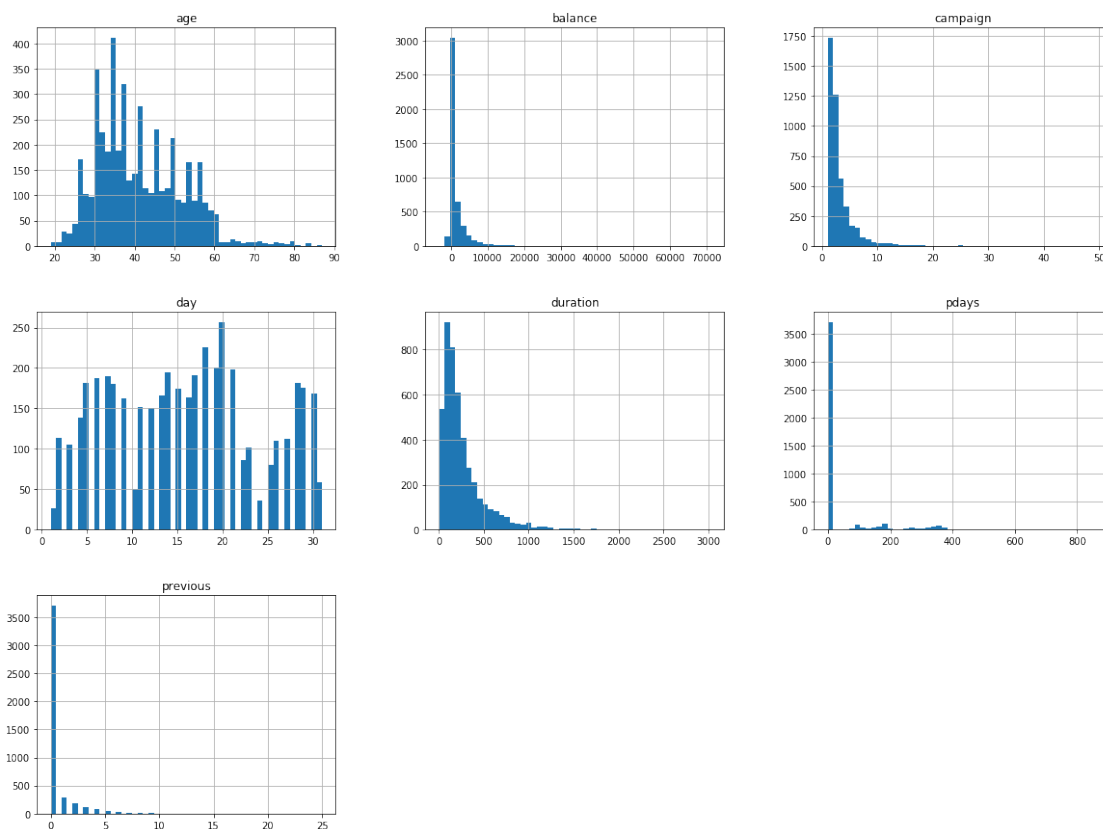
campaign      4521 non-null int64
pdays        4521 non-null int64
previous      4521 non-null int64
poutcome      4521 non-null object
response      4521 non-null object
dtypes: int64(7), object(10)
memory usage: 600.5+ KB

```

```

In [14]: %matplotlib inline
         bank.hist(bins=50, figsize=(20,15))
         plt.savefig('bank_hist.pdf')
         plt.show()

```



```

In [15]: # drop observations with missing data, if any
         bank.dropna()
         # examine the shape of input data after dropping missing data
         print(bank.shape)

```

```

(4521, 17)

```

```
In [16]: # look at the list of column names, note that y is the response
list(bank.columns.values)
```

```
Out[16]: ['age',
          'job',
          'marital',
          'education',
          'default',
          'balance',
          'housing',
          'loan',
          'contact',
          'day',
          'month',
          'duration',
          'campaign',
          'pdays',
          'previous',
          'poutcome',
          'response']
```

```
In [17]: # look at the beginning of the DataFrame
bank.head()
```

```
Out[17]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	30	unemployed	married	primary	no	1787	no	no	
1	33	services	married	secondary	no	4789	yes	yes	
2	35	management	single	tertiary	no	1350	yes	no	
3	30	management	married	tertiary	no	1476	yes	yes	
4	59	blue-collar	married	secondary	no	0	yes	no	

	contact	day	month	duration	campaign	pdays	previous	poutcome	response
0	cellular	19	oct	79	1	-1	0	unknown	no
1	cellular	11	may	220	1	339	4	failure	no
2	cellular	16	apr	185	1	330	1	failure	no
3	unknown	3	jun	199	4	-1	0	unknown	no
4	unknown	5	may	226	1	-1	0	unknown	no

```
In [18]: # mapping function to convert text no/yes to integer 0/1
convert_to_binary = {'no' : 0, 'yes' : 1}
```

```
# define binary variable for having credit in default
default = bank['default'].map(convert_to_binary)
```

```
# define binary variable for having a mortgage or housing loan
housing = bank['housing'].map(convert_to_binary)
```

```
# define binary variable for having a personal loan
loan = bank['loan'].map(convert_to_binary)
```

```
# define response variable to use in the model
response = bank['response'].map(convert_to_binary)
```

```
In [19]: # gather three explanatory variables and response into a numpy array
# here we use .T to obtain the transpose for the structure we want
model_data = np.array([np.array(default), np.array(housing), np.array(loan),
                        np.array(response)]).T

# examine the shape of model_data, which we will use in subsequent modeling
print(model_data.shape)

# the rest of the program should set up the modeling methods
# and evaluation within a cross-validation design
```

```
(4521, 4)
```



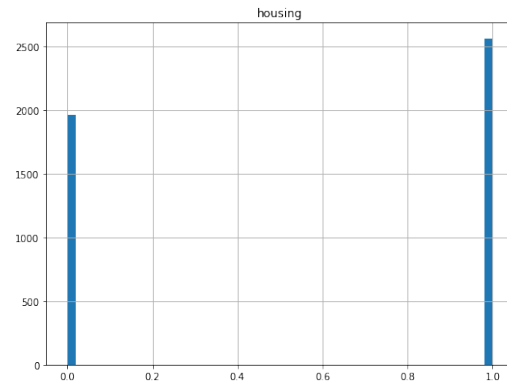
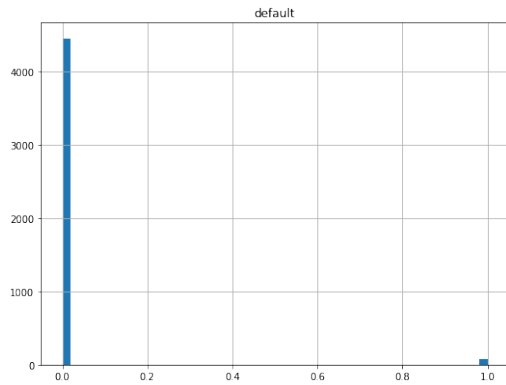
```
In [20]: model_data_df = pd.DataFrame(model_data)
model_data_df.columns = ['default', 'housing', 'loan', 'response']
```

```
In [21]: model_data_df.head()
```

```
Out[21]:
```

	default	housing	loan	response
0	0	0	0	0
1	0	1	1	0
2	0	1	0	0
3	0	1	1	0
4	0	1	0	0

```
In [22]: model_data_df.hist(bins=50, figsize=(20,15))
plt.savefig('model_data_df_hist.pdf')
plt.show()
```

```
In [23]: # value counts of the response variable
# there is a low rate of subscription to a term deposit
# only 521 clients have a term deposit
model_data_df.response.value_counts()
```

```
Out[23]: 0    4000
         1     521
         Name: response, dtype: int64
```

```
In [24]: #split data and response
```

```
model_data_df_X = model_data_df.drop('response', axis=1)
model_data_df_y = model_data_df.response.copy()
```

```
In [25]: model_data_df_X.head()
```

```
Out[25]:   default  housing  loan
0         0         0      0
1         0         1      1
2         0         1      0
3         0         1      1
4         0         1      0
```

```
In [26]: model_data_df_y.head()
```

```
Out[26]: 0    0
         1    0
         2    0
         3    0
         4    0
         Name: response, dtype: int64
```

0.3 Split the data into train and test sets

```
In [27]: #from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    model_data_df_X, model_data_df_y, test_size=0.2, random_state=42)
```

0.4 Logistic Regression Model, C = 100

```
In [28]: # Create a logistic regression model on the data
         # C is the hyperparameter controlling the regularization
         # strength of a Scikit-Learn LogisticRegression model.
         # The higher the value of C, the less the model is regularized.

         # Instantiate model
log_reg100 = LogisticRegression(C=100)

         # Fit the model
log_reg100.fit(X_train, y_train)

Out[28]: LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
        intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
        penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
        verbose=0, warm_start=False)

In [29]: # make predictions for the testing set
y_predictions100 = log_reg100.predict(X_test)

In [30]: # look at the first 10 entries for predictions
y_predictions100[0:10]

Out[30]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

In [31]: # look at the first 10 entries for the true values
y_test[0:10]

Out[31]: 2398    0
         800    0
         2288   0
         2344   0
```

```

3615    0
3548    0
1115    0
4053    0
838     0
4141    0
Name: response, dtype: int64

```

```

In [32]: # make predicted probabilities for the predictions
y_predict_prob100 = log_reg100.predict_proba(X_test)

```

```

In [33]: # Look at the first ten rows of predicted probabilities
# of response class membership.
# The first column is the predicted probability that the
# observation is a member of class 0.
# The second column is the predicted probability that the
# observation is a member of class 1.

```

```

y_predict_prob100[0:10]

```

```

Out[33]: array([[0.871, 0.129],
                [0.831, 0.169],
                [0.831, 0.169],
                [0.909, 0.091],
                [0.831, 0.169],
                [0.831, 0.169],
                [0.909, 0.091],
                [0.831, 0.169],
                [0.831, 0.169],
                [0.831, 0.169]])

```

```

In [34]: # the first argument is true values,
# the second argument is predicted values
# this produces a 2x2 numpy array (matrix)

```

```

confusion100 = metrics.confusion_matrix(y_test, y_predictions100)
print(confusion100)

```

```

[[807  0]
 [ 98  0]]

```

```

In [35]: # precision, recall, F1 score, and count of response variable
# of logistic regression model, C =100 when threshold set to 0.5

```

```

print(classification_report(y_test, y_predictions100))

```

```

precision    recall  f1-score   support

```

0	0.89	1.00	0.94	807
1	0.00	0.00	0.00	98
avg / total	0.80	0.89	0.84	905

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: U
'precision', 'predicted', average, warn_for)
```

0.5 Adjusting the classification threshold

```
In [36]: # store the predicted probabilities for class 1 of the response
# for the logistic regression, C = 100 model
y_pred_prob100_class1 = log_reg100.predict_proba(X_test)[: ,1]
```

```
In [37]: # look at the first ten entries for predicted
# probabilities for class 1
y_pred_prob100_class1[0:10]
```

```
Out[37]: array([0.129, 0.169, 0.169, 0.091, 0.169, 0.169, 0.091, 0.169, 0.169,
0.169])
```

```
In [38]: # the default threshold for predicted probabilities to be classified
# as 0 or 1 is 0.5
# let's see what would happen if the default is set to 0.1
# predict response if the predicted probability is greater than 0.1
# it will return 1 for all values above 0.1 and 0 otherwise
# results are 2D so we slice out the first column

y_pred_class100 = binarize(y_pred_prob100_class1.reshape(-1, 1), 0.1)
```

```
In [39]: # print the first 10 predicted classes with the lower threshold
y_pred_class100[0:10]
```

```
Out[39]: array([[1.],
[1.],
[1.],
[0.],
[1.],
[1.],
[0.],
[1.],
[1.],
[1.]])
```

```
In [40]: # previous confusion matrix (default threshold of 0.5)
print(confusion100)
```

```
[[807  0]
 [ 98  0]]
```

```
In [41]: # new confusion matrix (threshold set to 0.1)
```

```
confusion100_threshold1 = metrics.confusion_matrix(y_test, y_pred_class100)
print(confusion100_threshold1)
```

```
[[528 279]
 [ 50  48]]
```

```
In [42]: # precision, recall, F1 score, and count of response variable
# of logistic regression model, C =100 when threshold set to 0.1
```

```
print(classification_report(y_test, y_pred_class100))
```

	precision	recall	f1-score	support
0	0.91	0.65	0.76	807
1	0.15	0.49	0.23	98
avg / total	0.83	0.64	0.70	905

```
In [43]: # evaluate logistic regression model, C = 100
# when threshold set to 0.1
```

```
model_metrics(y_test, y_pred_class100)
```

The model metrics are:

accuracy: 0.63646408839779

classification error: 0.36353591160221

sensitivity: 0.4897959183673469

specificity: 0.654275092936803

false positive rate: 0.34572490706319703

precision: 0.14678899082568808

F1 score: 0.2258823529411765

confusion matrix:

```
Out[43]: array([[528, 279],
                [ 50,  48]])
```

```
In [44]: # evaluate logistic regression model, C = 100
# when threshold set to default 0.5
```

```
model_metrics(y_test, y_predictions100)
```

```
#y_predictions100 is the predicted response value
```

The model metrics are:

accuracy: 0.8917127071823204

classification error: 0.10828729281767957

sensitivity: 0.0

specificity: 1.0

false positive rate: 0.0

precision: nan

F1 score: 0.0

confusion matrix:

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:16: RuntimeWarning:
  app.launch_new_instance()
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: U
  'precision', 'predicted', average, warn_for)
```

```
Out[44]: array([[807,   0],
               [ 98,   0]])
```

0.6 ROC Curve

```
In [45]: # make predicted probabilities for the predictions
         # logistic regression, C = 100
         # store the predicted probabilities for class 1 of response
```

```
y_pred_prob100_class1 = log_reg100.predict_proba(X_test)[:, 1]
```

```
In [46]: # the first argument is true values,
         # the second argument is predicted probabilities

         # pass y_test and y_pred_prob
         # do not use y_pred_class, because it will
         # give incorrect results without generating an error
         # roc_curve returns 3 objects fpr, tpr, thresholds
         # fpr: false positive rate
         # tpr: true positive rate
```

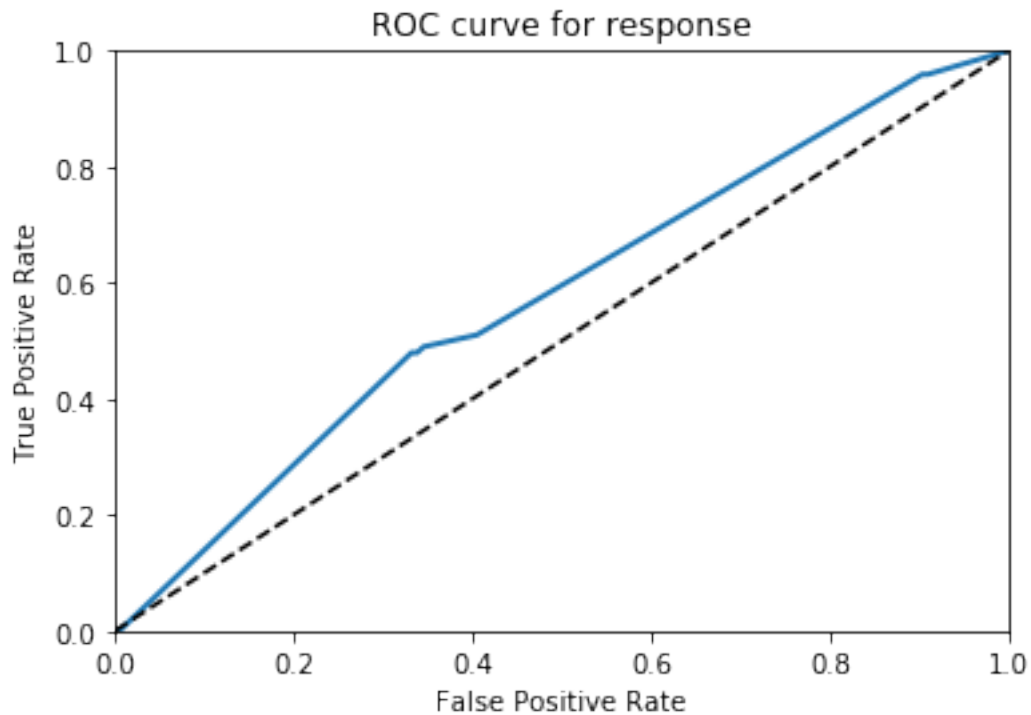
```
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob100_class1)
```

```
In [47]: # a check of thresholds
         thresholds
```

```
Out[47]: array([1.229, 0.229, 0.169, 0.129, 0.128, 0.092, 0.091, 0.068, 0.048])
```

```
In [48]: # plot the ROC curve for the logistic regression, C = 100
```

```
plot_roc_curve(fpr, tpr)
plt.savefig('ROC_logistic_C100_plot.pdf')
plt.show()
```



```
In [49]: # evaluate the logistic regression model, C = 100
# sensitivity and specificity when the threshold is set to 0.5
```

```
print('When the threshold is set to 0.5')
evaluate_threshold(0.5)
```

```
When the threshold is set to 0.5
Sensitivity: 0.0
Specificity: 1.0
```

```
In [50]: # evaluate the logistic regression model, C = 100
# sensitivity and specificity when the threshold is set to 0.1
```

```
print('When the threshold is set to 0.1')
evaluate_threshold(0.1)
```

```
When the threshold is set to 0.1
Sensitivity: 0.4897959183673469
```

Specificity: 0.654275092936803

```
In [51]: # AUC is the percentage of the ROC plot that is underneath the curve
# first argument is true values, second argument is predicted probabilities
# AUC for logistic regression, C = 100
```

```
roc_auc_y100_class1 = metrics.roc_auc_score(y_test, y_pred_prob100_class1)

print('\nThe AUC is:', roc_auc_y100_class1)
```

The AUC is: 0.5752977771034065

0.7 PR Curve

```
In [52]: # Given the imbalance of 0 to 1 in the response category
# let's see what the Precision-Recall curve looks like
# for the logistic regression model, C = 100
```

```
precision, recall, thresholds = precision_recall_curve(y_test, y_pred_prob100_class1)
```

```
In [53]: y_test[:10]
```

```
Out [53]: 2398    0
          800    0
          2288   0
          2344   0
          3615   0
          3548   0
          1115   0
          4053   0
          838    0
          4141   0
          Name: response, dtype: int64
```

```
In [54]: type(y_test)
```

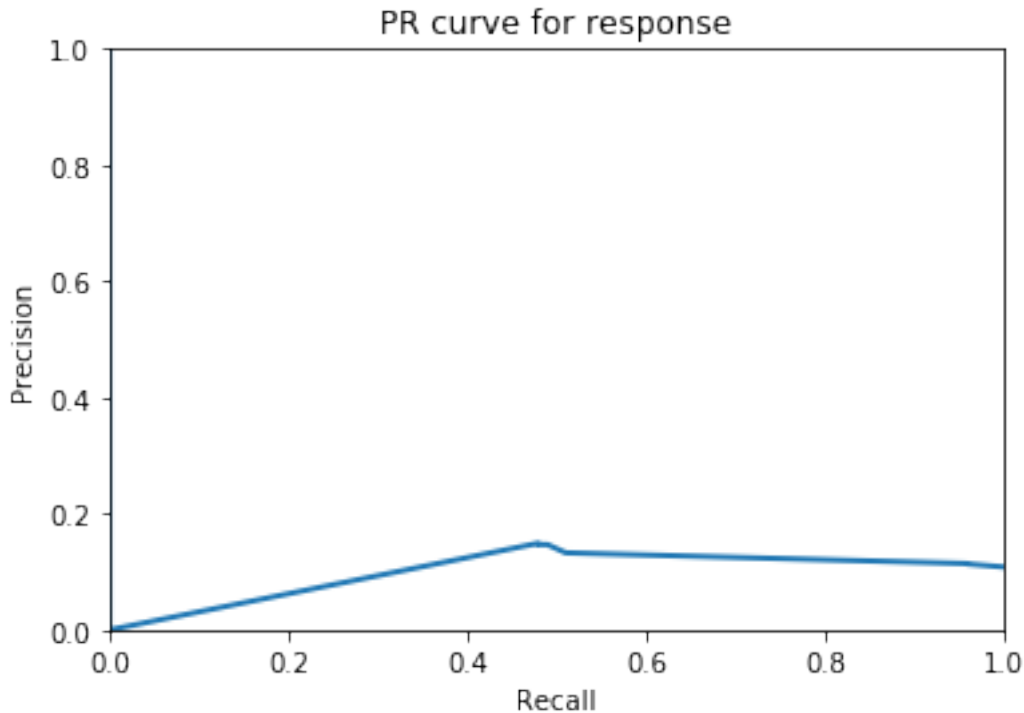
```
Out [54]: pandas.core.series.Series
```

```
In [55]: y_pred_prob100_class1[:10]
```

```
Out [55]: array([0.129, 0.169, 0.169, 0.091, 0.169, 0.169, 0.091, 0.169, 0.169,
                0.169])
```

```
In [56]: # plot the PR curve for logistic regression, C = 100
```

```
plot_pr_curve(recall, precision)
plt.savefig('PR_logistic_C100_plot.pdf')
plt.show()
```

0.8 Logistic Regression Model, C = 1000

```
In [57]: # Create a logistic regression model on the data
# C is the hyperparameter controlling the regularization
# strength of a Scikit-Learn LogisticRegression model.
# The higher the value of C, the less the model is regularized.
# Let's see if there is a difference with C = 1000
```

```
# Instantiate model
log_reg1000 = LogisticRegression(C=1000)

# Fit the model
log_reg1000.fit(X_train, y_train)
```

```
Out[57]: LogisticRegression(C=1000, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)
```

```
In [58]: # make predictions for the testing set
y_predictions1000 = log_reg1000.predict(X_test)
```

```
In [59]: # make predicted probabilities for the predictions
y_predict_proba1000 = log_reg1000.predict_proba(X_test)
```

```
In [60]: # store the predicted probabilities for class 1 of the response
# for the logistic regression, C = 1000 model

y_pred_prob1000_class1 = log_reg1000.predict_proba(X_test)[:,-1]

In [61]: # predict response if the predicted probability is greater than 0.1
# it will return 1 for all values above 0.1 and 0 otherwise
# results are 2D so we slice out the first column
y_pred_class1000 = binarize(y_pred_prob1000_class1.reshape(-1, 1), 0.1)

In [62]: # precision, recall, F1 score, and count of response variable
# of logistic regression model, C =1000 when threshold set to 0.5
```

```
print(classification_report(y_test, y_predictions1000))
```

	precision	recall	f1-score	support
0	0.89	1.00	0.94	807
1	0.00	0.00	0.00	98
avg / total	0.80	0.89	0.84	905

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UserWarning:
'precision', 'predicted', average, warn_for)
```

```
In [63]: # evaluate logistic regression model, C = 1000
# when threshold set to default 0.5
# model_metrics(y_known, y_pred):
```

```
model_metrics(y_test, y_predictions1000)
```

The model metrics are:

```
accuracy: 0.8917127071823204
classification error: 0.10828729281767957
sensitivity: 0.0
specificity: 1.0
false positive rate: 0.0
precision: nan
F1 score: 0.0
confusion matrix:
```


```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:16: RuntimeWarning:
app.launch_new_instance()
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UserWarning:
'precision', 'predicted', average, warn_for)
```

```
Out [63]: array([[807,  0],
                [ 98,  0]])
```

```
In [64]: # the first argument is true values,
# the second argument is predicted probabilities

# pass y_test and y_pred_prob
# do not use y_pred_class, because it will give
# incorrect results without generating an error
# roc_curve returns 3 objects fpr, tpr, thresholds
# fpr: false positive rate
# tpr: true positive rate

fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob1000_class1)
plot_roc_curve(fpr, tpr)
plt.savefig('ROC_logistic_C1000_plot.pdf')
plt.show()
```



```
In [65]: # evaluate the logistic regression model, C = 1000
# sensitivity and specificity when the threshold is set to 0.5

print('When the threshold is set to 0.5')
evaluate_threshold(0.5)
```

When the threshold is set to 0.5
Sensitivity: 0.0
Specificity: 1.0

```
In [66]: # evaluate the logistic regression model, C = 1000
         # sensitivity and specificity when the threshold is set to 0.1

         print('When the threshold is set to 0.1')
         evaluate_threshold(0.1)
```

When the threshold is set to 0.1
Sensitivity: 0.4897959183673469
Specificity: 0.654275092936803

```
In [67]: # AUC is the percentage of the ROC plot that is underneath the curve
         # first argument is true values, second argument is predicted probabilities
         # AUC for logistic regression, C = 1000


         roc_auc_y100_class1 = metrics.roc_auc_score(y_test, y_pred_prob1000_class1)
         print('\nThe AUC is:', roc_auc_y100_class1)
```

The AUC is: 0.5752977771034065

```
In [68]: # Given the imbalance of 0 to 1 in the response category
         # let's see what the Precision-Recall curve looks like
         # for the logistic regression model, C = 1000

         precision, recall, thresholds = precision_recall_curve(y_test, y_pred_prob1000_class1)

         # plot the PR curve for logistic regression, C = 1000
         plot_pr_curve(recall, precision)
         plt.savefig('PR_logistic_C1000_plot.pdf')
         plt.show()
```



0.9 Naive Bayes Classifier

In [69]: *# Create a Naive Bayes Classifier on the data*

```
# Instantiate model
clf = BernoulliNB()

# Fit the model
clf.fit(X_train, y_train)
```

Out [69]: BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)

In [70]: *# make predictions for the testing set*
clf_y_predictions = clf.predict(X_test)

In [71]: *# make predicted probabilities for the predictions*
clf_y_predict_proba = clf.predict_proba(X_test)

In [72]: *# store the predicted probabilities for class 1 of the response*
for the naive bayes classifier

```
clf_y_pred_proba_class1 = clf.predict_proba(X_test)[: ,1]
```

```
In [73]: # precision, recall, F1 score, and count of response variable
         # of naive bayes classifier
```

```
print(classification_report(y_test, clf_y_predictions))
```

	precision	recall	f1-score	support
0	0.89	1.00	0.94	807
1	0.00	0.00	0.00	98
avg / total	0.80	0.89	0.84	905

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: U
'precision', 'predicted', average, warn_for)
```

```
In [74]: # evaluate naive bayes classifier
         # when threshold set to default 0.5
         # model_metrics(y_known, y_pred):
```

```
model_metrics(y_test, clf_y_predictions)
```

The model metrics are:

accuracy: 0.8917127071823204

classification error: 0.10828729281767957

sensitivity: 0.0

specificity: 1.0

false positive rate: 0.0

precision: nan

F1 score: 0.0

confusion matrix:

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:16: RuntimeWarning:
app.launch_new_instance()
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: U
'precision', 'predicted', average, warn_for)
```

```
Out[74]: array([[807,  0],
               [ 98,  0]])
```

```
In [75]: # first argument is true values,
         # second argument is predicted probabilities


         # pass y_test and y_pred_prob
         # do not use y_pred_class, because it will give
```

```

# incorrect results without generating an error
# roc_curve returns 3 objects fpr, tpr, thresholds
# fpr: false positive rate
# tpr: true positive rate

fpr, tpr, thresholds = roc_curve(y_test, clf_y_pred_prob_class1)
plot_roc_curve(fpr, tpr)
plt.savefig('ROC_bernoulli_plot.pdf')
plt.show()

```



```

In [76]: # evaluate the naive bayes classifier
# sensitivity and specificity when the threshold is set to 0.5

print('When the threshold is set to 0.5')
evaluate_threshold(0.5)

```

```

When the threshold is set to 0.5
Sensitivity: 0.0
Specificity: 1.0

```

```

In [77]: # evaluate the naive bayes classifier
# sensitivity and specificity when the threshold is set to 0.1

```

```
print('When the threshold is set to 0.1')
evaluate_threshold(0.1)
```

When the threshold is set to 0.1

Sensitivity: 0.4897959183673469

Specificity: 0.654275092936803


```
In [78]: # AUC is the percentage of the ROC plot that is underneath the curve
# first argument is true values, second argument is predicted probabilities
# AUC for naive bayes classifier
```

```
roc_auc_y100_class1 = metrics.roc_auc_score(y_test, clf_y_pred_prob_class1)
print('\nThe AUC is:', roc_auc_y100_class1)
```

The AUC is: 0.5752977771034065

```
In [79]: # Given the imbalance of 0 to 1 in the response category
# let's see what the Precision-Recall curve looks like
# for the naive bayes classifier
```

```
precision, recall, thresholds = precision_recall_curve(y_test, clf_y_pred_prob_class1)
plot_pr_curve(recall, precision)
plt.savefig('PR_bernoulli_plot.pdf')
plt.show()
```



0.10 Cross Validation of Logistic Regression Model, C = 100

```
In [80]: # calculate cross-validated AUC for logistic regression model
```

```
#log_reg100 = LogisticRegression(C=100)
print('Cross validated AUC for Logistic Regression:')
cross_log_auc = cross_val_score(log_reg100, X_train, y_train, cv=10, scoring='roc_auc')

display_scores(cross_log_auc)
```

Cross validated AUC for Logistic Regression:

Scores: [0.627 0.606 0.555 0.68 0.632 0.592 0.554 0.599 0.577 0.622]

Mean: 0.6045377265827597

Standard deviation: 0.03632462069250767

0.11 Cross Validation of Naive Bayes Classification

```
In [81]: # calculate cross-validated AUC for naive Bayes classification
```

```
print('Cross validated AUC for Naive Bayes Classification:')
cross_clf_auc = cross_val_score(clf, X_train, y_train, cv=10, scoring='roc_auc')

display_scores(cross_clf_auc)
```

Cross validated AUC for Naive Bayes Classification:

Scores: [0.627 0.606 0.547 0.68 0.632 0.592 0.554 0.599 0.577 0.622]

Mean: 0.6036947033269457

Standard deviation: 0.03753998048084224

0.12 Interpreting the Logistic Regression Model, C = 100

```
In [82]: log_reg100.coef_
```

```
Out[82]: array([[ 0.382, -0.708, -0.695]])
```

```
In [83]: model_data_df_X.head()
```

```
Out[83]:
```

	default	housing	loan
0	0	0	0
1	0	1	1
2	0	1	0
3	0	1	1
4	0	1	0

```

In [84]: #yes = 1, no = 0
         #response = has the client subscribed to a term deposit?
         #for credit in default = yes, the log of the odds of response increase by 0.382
         #for housing loan = yes, the log of the odds of response decrease by -0.708
         #for personal loan = yes, the log of the odds of response decrease by -0.695

In [85]: clf.coef_

Out[85]: array([[ -3.855,  -0.905,  -2.415]])

In [86]: math.exp(-3.855)

         #for client with credit in default, there is a 2% increase in having a term deposit

Out[86]: 0.02117360331011653

In [87]: math.exp(-0.905)

         #for client with housing loan, there is a 40% decrease in having a term deposit

Out[87]: 0.4045418851030188

In [88]: math.exp(-2.415)

         #for client with personal loan, there is a 8.9% decrease in having a term deposit

Out[88]: 0.08936733892175319

```