

Wanat_Assignment_1_Python

June 30, 2019

```
In [1]: # Jump-Start Example: Python analysis of MSPA Software Survey

# Update 2017-09-21 by Tom Miller and Kelsey O'Neill
# Update 2018-06-30 by Tom Miller v005 transformation code added

# tested under Python 3.6.1 :: Anaconda custom (x86_64)
# on Windows 10.0 and Mac OS Sierra 10.12.2

# shows how to read in data from a comma-delimited text file
# manipulate data, create new count variables, define categorical variables,
# work with dictionaries and lambda mapping functions for recoding data

# visualizations in this program are routed to external pdf files
# so they may be included in printed or electronic reports

# prepare for Python version 3x features and functions
# these two lines of code are needed for Python 2.7 only
# commented out for Python 3.x versions
# from __future__ import division, print_function
# from future_builtins import ascii, filter, hex, map, oct, zip

In [2]: # external libraries for visualizations and data manipulation
# ensure that these packages have been installed prior to calls
import pandas as pd # data frame operations
import numpy as np # arrays and math functions
import pandas_profiling
import matplotlib
import matplotlib.pyplot as plt # static plotting
import seaborn as sns # pretty plotting, including heat map

In [3]: # correlation heat map setup for seaborn
def corr_chart(df_corr):
    corr=df_corr.corr()
    #screen top half to get a triangle
    top = np.zeros_like(corr, dtype=np.bool)
    top[np.triu_indices_from(top)] = True
    fig=plt.figure()
```

```

fig, ax = plt.subplots(figsize=(12,12))
sns.heatmap(corr, mask=top, cmap='coolwarm',
            center = 0, square=True,
            linewidths=.5, cbar_kws={'shrink':.5},
            annot = True, annot_kws={'size': 9}, fmt = '.3f')
plt.xticks(rotation=45) # rotate variable labels on columns (x axis)
plt.yticks(rotation=0) # use horizontal variable labels on rows (y axis)
plt.title('Correlation Heat Map')
plt.savefig('plot-corr-map.pdf',
            bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
            orientation='portrait', papertype=None, format=None,
            transparent=True, pad_inches=0.25, frameon=None)

```

```
np.set_printoptions(precision=3)
```

```

In [4]: # read in comma-delimited text file, creating a pandas DataFrame object
# note that IPAddress is formatted as an actual IP address
# but is actually a random-hash of the original IP address
valid_survey_input = pd.read_csv('mspa-survey-data.csv')

# use the RespondentID as label for the rows... the index of DataFrame
valid_survey_input.set_index('RespondentID', drop = True, inplace = True)

# examine the structure of the DataFrame object
print('\nContents of initial survey data -----')

# could use len() or first index of shape() to get number of rows/observations
print('\nNumber of Respondents =', len(valid_survey_input))

```

Contents of initial survey data -----

Number of Respondents = 207

```

In [5]: #The shape attribute tells us the numbers of rows and columns in the data frame
print('\nThe shape of the dataframe (rows, columns):')
valid_survey_input.shape

```

The shape of the dataframe (rows, columns):

Out[5]: (207, 40)

```

In [6]: # provide the number of columns in data frame
print('\nNumber of columns in data frame = ', valid_survey_input.shape[1])

```

Number of columns in data frame = 40

```
In [7]: # show the column/variable names of the DataFrame
        # note that RespondentID is no longer present
        print(valid_survey_input.columns)
```

```
Index(['Personal_JavaScalaSpark', 'Personal_JavaScriptHTMLCSS',
      'Personal_Python', 'Personal_R', 'Personal_SAS',
      'Professional_JavaScalaSpark', 'Professional_JavaScriptHTMLCSS',
      'Professional_Python', 'Professional_R', 'Professional_SAS',
      'Industry_JavaScalaSpark', 'Industry_JavaScriptHTMLCSS',
      'Industry_Python', 'Industry_R', 'Industry_SAS',
      'Python_Course_Interest', 'Foundations_DE_Course_Interest',
      'Analytics_App_Course_Interest', 'Systems_Analysis_Course_Interest',
      'Courses_Completed', 'PREDICT400', 'PREDICT401', 'PREDICT410',
      'PREDICT411', 'PREDICT413', 'PREDICT420', 'PREDICT422', 'PREDICT450',
      'PREDICT451', 'PREDICT452', 'PREDICT453', 'PREDICT454', 'PREDICT455',
      'PREDICT456', 'PREDICT457', 'OtherPython', 'OtherR', 'OtherSAS',
      'Other', 'Graduate_Date'],
      dtype='object')
```

```
In [8]: # examine the data types for each column in the data frame
        print(valid_survey_input.dtypes)
```

Personal_JavaScalaSpark	int64
Personal_JavaScriptHTMLCSS	int64
Personal_Python	int64
Personal_R	int64
Personal_SAS	int64
Professional_JavaScalaSpark	int64
Professional_JavaScriptHTMLCSS	int64
Professional_Python	int64
Professional_R	int64
Professional_SAS	int64
Industry_JavaScalaSpark	int64
Industry_JavaScriptHTMLCSS	int64
Industry_Python	int64
Industry_R	int64
Industry_SAS	int64
Python_Course_Interest	float64
Foundations_DE_Course_Interest	float64
Analytics_App_Course_Interest	float64
Systems_Analysis_Course_Interest	float64
Courses_Completed	float64
PREDICT400	object
PREDICT401	object
PREDICT410	object
PREDICT411	object
PREDICT413	object

```

PREDICT420          object
PREDICT422          object
PREDICT450          object
PREDICT451          object
PREDICT452          object
PREDICT453          object
PREDICT454          object
PREDICT455          object
PREDICT456          object
PREDICT457          object
OtherPython         object
OtherR              object
OtherSAS            object
Other               object
Graduate_Date       object
dtype: object

```

```

In [9]: #total number of NaN values in each column
        valid_survey_input.isnull().sum()

```

```

Out[9]: Personal_JavaScalaSpark          0
        Personal_JavaScriptHTMLCSS       0
        Personal_Python                   0
        Personal_R                        0
        Personal_SAS                      0
        Professional_JavaScalaSpark       0
        Professional_JavaScriptHTMLCSS    0
        Professional_Python               0
        Professional_R                    0
        Professional_SAS                  0
        Industry_JavaScalaSpark           0
        Industry_JavaScriptHTMLCSS        0
        Industry_Python                   0
        Industry_R                        0
        Industry_SAS                      0
        Python_Course_Interest            1
        Foundations_DE_Course_Interest    7
        Analytics_App_Course_Interest     4
        Systems_Analysis_Course_Interest  7
        Courses_Completed                 20
        PREDICT400                        44
        PREDICT401                        36
        PREDICT410                        62
        PREDICT411                        94
        PREDICT413                       148
        PREDICT420                        80
        PREDICT422                       159

```

PREDICT450	190
PREDICT451	200
PREDICT452	194
PREDICT453	196
PREDICT454	202
PREDICT455	177
PREDICT456	201
PREDICT457	203
OtherPython	202
OtherR	193
OtherSAS	205
Other	181
Graduate_Date	3
dtype:	int64

0.1 Profile Analysis of Data Set

In [10]: *#The ProfileReport in the pandas_profiling package provides a summary of the dataset*

```
pandas_profiling.ProfileReport(valid_survey_input)
```

Out [10]: <pandas_profiling.ProfileReport at 0x113ad55f8>

In [11]: *# Save the pandas profile report to html file*

```
profile = pandas_profiling.ProfileReport(valid_survey_input)
profile.to_file()
```

0.2 Analysis of Course Completion

In [12]: *# Analysis of Course Completion*

shorten the variable/column names for software preference variables

```
survey_df = valid_survey_input.rename(index=str, columns={
    'Personal_JavaScalaSpark': 'My_Java',
    'Personal_JavaScriptHTMLCSS': 'My_JS',
    'Personal_Python': 'My_Python',
    'Personal_R': 'My_R',
    'Personal_SAS': 'My_SAS',
    'Professional_JavaScalaSpark': 'Prof_Java',
    'Professional_JavaScriptHTMLCSS': 'Prof_JS',
    'Professional_Python': 'Prof_Python',
    'Professional_R': 'Prof_R',
    'Professional_SAS': 'Prof_SAS',
    'Industry_JavaScalaSpark': 'Ind_Java',
    'Industry_JavaScriptHTMLCSS': 'Ind_JS',
    'Industry_Python': 'Ind_Python',
    'Industry_R': 'Ind_R',
    'Industry_SAS': 'Ind_SAS'})
```

```
In [13]: # Print the first five rows of the survey_df that contains
# shortened variable/column names for software preference variables
# to check that it worked
```

```
print(pd.DataFrame.head(survey_df))
```

RespondentID	My_Java	My_JS	My_Python	My_R	My_SAS	Prof_Java	Prof_JS	\
5135740122	0	0	0	50	50	0	0	
5133300037	10	10	50	30	0	25	25	
5132253300	20	0	40	40	0	0	0	
5132096630	10	10	25	35	20	10	10	
5131990362	20	0	0	70	10	20	0	

RespondentID	Prof_Python	Prof_R	Prof_SAS	...	PREDICT453	\
5135740122	0	25	75	...	NaN	
5133300037	30	20	0	...	NaN	
5132253300	40	40	20	...	NaN	
5132096630	25	35	20	...	NaN	
5131990362	0	80	0	...	NaN	

RespondentID	PREDICT454	PREDICT455	PREDICT456	PREDICT457	OtherPython	\
5135740122	NaN	NaN	NaN	NaN	NaN	
5133300037	NaN	NaN	NaN	NaN	NaN	
5132253300	NaN	NaN	NaN	NaN	NaN	
5132096630	NaN	NaN	NaN	NaN	NaN	
5131990362	NaN	NaN	NaN	NaN	NaN	

RespondentID	OtherR	OtherSAS	Other	Graduate_Date
5135740122	NaN	NaN	NaN	NaN
5133300037	NaN	NaN	NaN	Spring 2018
5132253300	NaN	NaN	NaN	Fall 2018
5132096630	NaN	NaN	NaN	Fall 2017
5131990362	NaN	NaN	CS-435 with Weka	Fall 2018

```
[5 rows x 40 columns]
```

```
In [14]: # descriptive statistics for one variable
print('\nDescriptive statistics for courses completed -----')
print(survey_df['Courses_Completed'].describe())
```

```
Descriptive statistics for courses completed -----
count    187.000000
```

```
mean      6.342246
std       3.170849
min       1.000000
25%      4.000000
50%      6.000000
75%      9.000000
max      12.000000
Name: Courses_Completed, dtype: float64
```

```
In [15]: #Summarize courses completed by displaying the number of students in each category
courses_completed_counts = survey_df['Courses_Completed'].value_counts()
print('\nThe number of courses completed by students: ')
courses_completed_counts.sort_index()
```

The number of courses completed by students:

```
Out[15]: 1.0      6
        2.0     25
        3.0     13
        4.0     13
        5.0     24
        6.0     16
        7.0     24
        8.0     11
        9.0     14
       10.0     20
       11.0     11
       12.0     10
        Name: Courses_Completed, dtype: int64
```

```
In [16]: #Display the percentages of each course completed category
print('\nPercentage of the Number of Courses Completed: ')
((courses_completed_counts/sum(courses_completed_counts))*100)
```

Percentage of the Number of Courses Completed:

```
Out[16]: 2.0      13.368984
        5.0      12.834225
        7.0      12.834225
       10.0      10.695187
        6.0       8.556150
        9.0       7.486631
        3.0       6.951872
        4.0       6.951872
```

```

8.0      5.882353
11.0     5.882353
12.0     5.347594
1.0      3.208556
Name: Courses_Completed, dtype: float64

```

```

In [17]: # Display the percentages of each course completed category
# Another way to calculate
print('\nPercentage of the Number of Courses Completed: ')
survey_df['Courses_Completed'].value_counts(normalize=True)*100

```

Percentage of the Number of Courses Completed:

```

Out[17]: 2.0      13.368984
5.0      12.834225
7.0      12.834225
10.0     10.695187
6.0       8.556150
9.0       7.486631
3.0       6.951872
4.0       6.951872
8.0       5.882353
11.0      5.882353
12.0      5.347594
1.0       3.208556
Name: Courses_Completed, dtype: float64

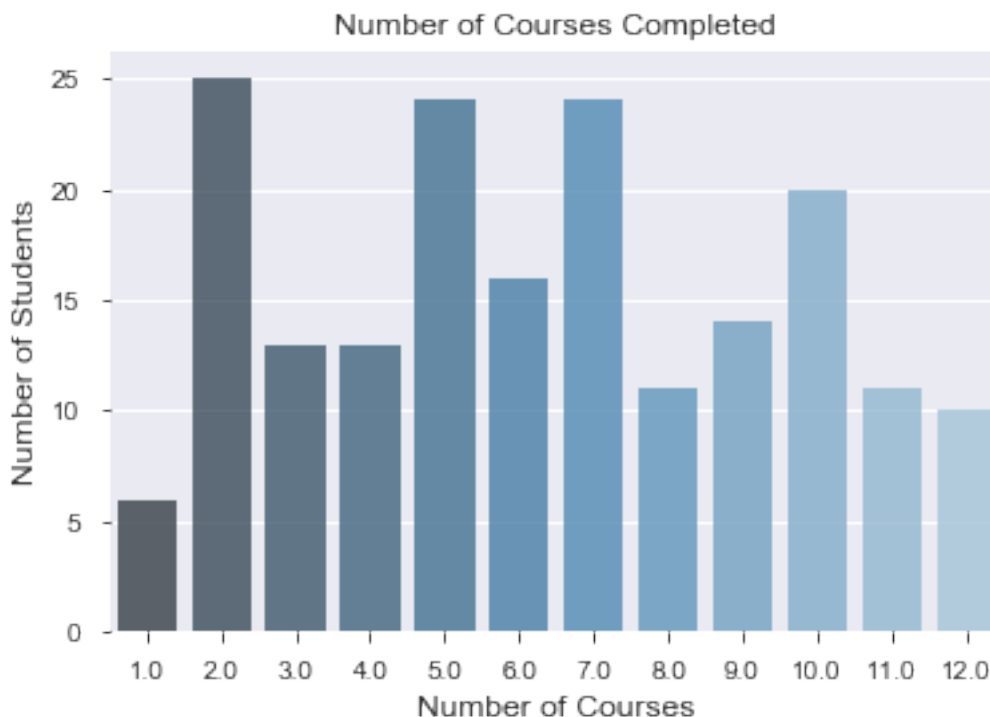
```

```

In [18]: # Examine the number of courses completed by student with a bar graph

%matplotlib inline
course_completed_date_fig, ax = plt.subplots()
sns.barplot(y = courses_completed_counts.values,
            x = courses_completed_counts.index, alpha=0.8,
            palette="Blues_d").set_title('Number of Courses Completed')
ax.set_xlabel('Number of Courses', fontsize=12)
ax.set_ylabel('Number of Students', fontsize=12)
course_completed_date_fig.savefig('CourseCompleted' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)

```

```
In [19]: # The selected columns were placed into a new dataframe.
# These columns are courses with a programming component
# that were completed by the respondent at the time of the survey.
# If the course was completed, the row contains a entry for that course.
# If not completed, the row contains a 'NaN'.

class_completed_df = valid_survey_input[['PREDICT400', 'PREDICT401', 'PREDICT410',
                                           'PREDICT411', 'PREDICT413', 'PREDICT420',
                                           'PREDICT422', 'PREDICT450', 'PREDICT451',
                                           'PREDICT452', 'PREDICT453', 'PREDICT454',
                                           'PREDICT455', 'PREDICT456', 'PREDICT457']]
```

```
In [20]: # View new dataframe for classes completed at time of survey
```

```
class_completed_df.head()
```

```
Out[20]:
```

RespondentID	PREDICT400	\
5135740122		NaN
5133300037	PREDICT 400 Math for Modelers (Python)	
5132253300	PREDICT 400 Math for Modelers (Python)	
5132096630	PREDICT 400 Math for Modelers (Python)	
5131990362		NaN

	PREDICT401 \
RespondentID	
5135740122	NaN
5133300037	PREDICT 401 Introduction to Statistical Analys...
5132253300	PREDICT 401 Introduction to Statistical Analys...
5132096630	PREDICT 401 Introduction to Statistical Analys...
5131990362	PREDICT 401 Introduction to Statistical Analys...

	PREDICT410 \
RespondentID	
5135740122	NaN
5133300037	PREDICT 410 Regression and Multivariate Analys...
5132253300	NaN
5132096630	PREDICT 410 Regression and Multivariate Analys...
5131990362	PREDICT 410 Regression and Multivariate Analys...

	PREDICT411 PREDICT413 \
RespondentID	
5135740122	NaN NaN
5133300037	PREDICT 411 Generalized Linear Models (SAS) NaN
5132253300	NaN NaN
5132096630	PREDICT 411 Generalized Linear Models (SAS) NaN
5131990362	PREDICT 411 Generalized Linear Models (SAS) NaN

	PREDICT420 PREDICT422 \
RespondentID	
5135740122	NaN NaN
5133300037	PREDICT 420 Database Systems and Data Preparat... NaN
5132253300	PREDICT 420 Database Systems and Data Preparat... NaN
5132096630	PREDICT 420 Database Systems and Data Preparat... NaN
5131990362	NaN NaN

	PREDICT450 PREDICT451 PREDICT452 PREDICT453 PREDICT454 \
RespondentID	
5135740122	NaN NaN NaN NaN NaN
5133300037	NaN NaN NaN NaN NaN
5132253300	NaN NaN NaN NaN NaN
5132096630	NaN NaN NaN NaN NaN
5131990362	NaN NaN NaN NaN NaN

	PREDICT455 PREDICT456 PREDICT457
RespondentID	
5135740122	NaN NaN NaN
5133300037	NaN NaN NaN
5132253300	NaN NaN NaN
5132096630	NaN NaN NaN
5131990362	NaN NaN NaN

```
In [21]: # It is known that the survey dataframe contains 207 rows.
# The class_completed_df will be examined for null values with the isnull function.
# The number of null values will be subtracted from 207, and the resulting number
# will be the number of respondents that have completed the course.

done_df = 207 - class_completed_df.isnull().sum()

print('\nThe number of respondents that have completed the following courses:')
done_df.sort_values(ascending=False)
```

The number of respondents that have completed the following courses:

```
Out [21]: PREDICT401    171
          PREDICT400    163
          PREDICT410    145
          PREDICT420    127
          PREDICT411    113
          PREDICT413     59
          PREDICT422     48
          PREDICT455     30
          PREDICT450     17
          PREDICT452     13
          PREDICT453     11
          PREDICT451      7
          PREDICT456      6
          PREDICT454      5
          PREDICT457      4
          dtype: int64
```

0.3 Analysis of Software Preference, Part I

```
In [22]: # Examine the columns for respondent desire to learn for each
# of the five language/software options.
# The describe function provides a basic statistical analysis.

survey_df.iloc[:,0:5].describe()
```

```
Out [22]:
```

	My_Java	My_JS	My_Python	My_R	My_SAS
count	207.000000	207.000000	207.000000	207.000000	207.000000
mean	10.135266	4.797101	31.304348	37.125604	16.637681
std	11.383477	6.757764	15.570982	14.576003	13.626400
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000	30.000000	5.000000
50%	9.000000	0.000000	30.000000	35.000000	15.000000
75%	20.000000	10.000000	40.000000	50.000000	25.000000
max	70.000000	30.000000	90.000000	100.000000	75.000000

```
In [23]: # Place the columns for respondent desire to learn for each
# of the five language/software options into a new data frame.
# The dataframe columns are 'melted' together in order to examine in
# a type of boxplot called a violinplot.
# The shape in a violinplot demonstrates the distribution of the data.


# The columns for respondent desire to learn each of the
# five language/software options were examined.

personal_df = survey_df.iloc[:,0:5]

personal_melted_df = personal_df.melt(var_name='groups', value_name='vals')

personal_desire_learn_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=personal_melted_df)
ax.set_title('Personal Desire to Learn Language or Software System', fontsize=18)
personal_desire_learn_fig.savefig('PersonalDesireToLearn' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

Personal Desire to Learn Language or Software System



```
In [24]: # Examine the columns for respondent professional need to learn
# for each of the five language/software options.
# The describe function provides a basic statistical analysis.

survey_df.iloc[:,5:10].describe()
```

```
Out [24]:
```

	Prof_Java	Prof_JS	Prof_Python	Prof_R	Prof_SAS
count	207.000000	207.000000	207.000000	207.000000	207.000000
mean	9.251208	5.840580	30.028986	36.415459	18.463768
std	13.167505	10.812555	19.144802	20.847606	18.831841
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000	25.000000	0.000000
50%	5.000000	0.000000	30.000000	33.000000	15.000000
75%	15.000000	10.000000	40.000000	50.000000	30.000000
max	80.000000	100.000000	100.000000	100.000000	100.000000

```
In [25]: # Place the columns for respondent professional need to learn for each
# of the five language/software options into a new data frame.
# The dataframe columns are 'melted' together in order to examine in
# a type of boxplot called a violinplot.
# The shape in a violinplot demonstrates the distribution of the data.

# The columns for respondent professional need to learn each of the
# five language/software options were examined.
professional_df = survey_df.iloc[:,5:10]

professional_melted_df = professional_df.melt(var_name='groups', value_name='vals')

professional_need_learn_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=professional_melted_df)
ax.set_title('Professional Need to Learn Language or Software System', fontsize=18)
professional_need_learn_fig.savefig('ProfessionalNeedToLearn' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

Professional Need to Learn Language or Software System



```
In [26]: # Examine the columns for respondent industry importance and prevalence to learn
# for each of the five language/software options.
# The describe function provides a basic statistical analysis.
```

```
survey_df.iloc[:,10:15].describe()
```

```
Out [26]:
```

	Ind_Java	Ind_JS	Ind_Python	Ind_R	Ind_SAS
count	207.000000	207.000000	207.000000	207.000000	207.000000
mean	11.942029	6.966184	29.772947	32.434783	18.884058
std	14.706399	10.030721	17.959816	15.912209	19.137623
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000	22.500000	0.000000
50%	5.000000	0.000000	30.000000	30.000000	15.000000
75%	20.000000	10.000000	40.000000	40.000000	30.000000
max	70.000000	50.000000	95.000000	85.000000	100.000000

```
In [27]: # Place the columns for respondent industry importance and prevalence to learn for ea
# of the five language/software options into a new data frame.
# The dataframe columns are 'melted' together in order to examine in
# a type of boxplot called a violinplot.
# The shape in a violinplot demonstrates the distribution of the data.
```

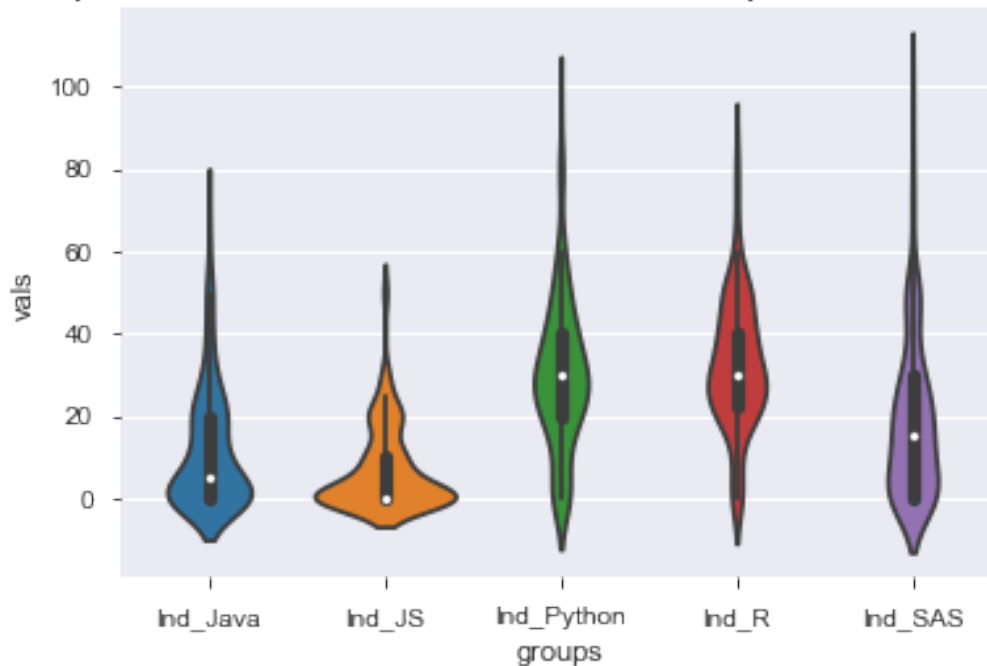
```
# The columns for respondent industry importance and prevalence to learn
# each of the five language/software options were examined.
```

```
industry_df = survey_df.iloc[:,10:15]
```

```
industry_melted_df = industry_df.melt(var_name='groups', value_name='vals')
```

```
industry_importance_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=industry_melted_df)
ax.set_title('Importance and Prevalence in Respondent Industry', fontsize=18)
industry_importance_fig.savefig('IndustryImportance' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```


Importance and Prevalence in Respondent Industry



In [28]: *# Examining the Java/Scala/Spark responses*


```
java_df = survey_df.iloc[:, [0, 5, 10]]
java_melted_df = java_df.melt(var_name='groups', value_name='vals')

java_scala_spark_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=java_melted_df)
ax.set_title('Java/Scala/Spark', fontsize=18)
java_scala_spark_fig.savefig('JavaScalaSpark' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```




```
In [29]: # Examining the JavaScript/HTML/CSS responses
js_df = survey_df.iloc[:, [1, 6, 11]]
js_melted_df = js_df.melt(var_name='groups', value_name='vals')

javascript_html_css_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=js_melted_df)
ax.set_title('JavaScript/HTML/CSS', fontsize=18)
javascript_html_css_fig.savefig('JavaScriptHTMLCSS' + '.pdf',
                                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                                orientation='portrait', papertype=None, format=None,
                                transparent=True, pad_inches=0.25, frameon=None)
```


```
In [30]: # Examining the Python responses
python_df = survey_df.iloc[:, [2, 7, 12]]
python_melted_df = python_df.melt(var_name='groups', value_name='vals')

python_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=python_melted_df)
ax.set_title('Python', fontsize=18)
python_fig.savefig('Python' + '.pdf',
                    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                    orientation='portrait', papertype=None, format=None,
                    transparent=True, pad_inches=0.25, frameon=None)
```



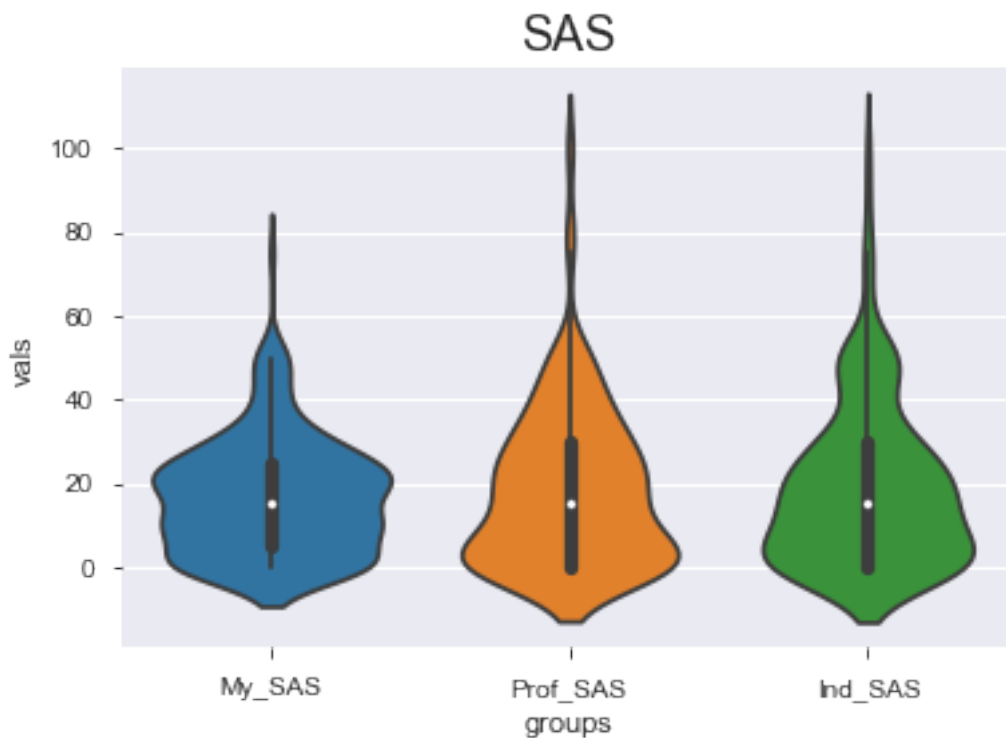
```
In [31]: # Examining the R responses
r_df = survey_df.iloc[:, [3, 8, 13]]
r_melted_df = r_df.melt(var_name='groups', value_name='vals')

rlanguage_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=r_melted_df)
ax.set_title('R', fontsize=18)
rlanguage_fig.savefig('Rlanguage' + '.pdf',
                      bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                      orientation='portrait', papertype=None, format=None,
                      transparent=True, pad_inches=0.25, frameon=None)
```



```
In [32]: # Examining the SAS responses
sas_df = survey_df.iloc[:, [4, 9, 14]]
sas_melted_df = sas_df.melt(var_name='groups', value_name='vals')

sas_fig, ax = plt.subplots()
sns.violinplot(x="groups", y="vals", data=sas_melted_df)
ax.set_title('SAS', fontsize=18)
sas_fig.savefig('SAS' + '.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)
```



0.4 Analysis of Software Preference, Part II

In [33]: *# define subset of survey DataFrame for analysis of software preferences*
`software_df = survey_df.loc[:, 'My_Java':'Ind_SAS']`

In [34]: *# display the first five lines of the data frame that contains*
personal, professional, and industry preferences by language and software
`display(pd.DataFrame.head(software_df))`

	My_Java	My_JS	My_Python	My_R	My_SAS	Prof_Java	Prof_JS	\
RespondentID								
5135740122	0	0	0	50	50	0	0	
5133300037	10	10	50	30	0	25	25	
5132253300	20	0	40	40	0	0	0	
5132096630	10	10	25	35	20	10	10	
5131990362	20	0	0	70	10	20	0	

	Prof_Python	Prof_R	Prof_SAS	Ind_Java	Ind_JS	Ind_Python	\
RespondentID							
5135740122	0	25	75	0	0	0	
5133300037	30	20	0	20	25	40	
5132253300	40	40	20	30	0	30	
5132096630	25	35	20	10	10	25	

5131990362	0	80	0	40	0	0
------------	---	----	---	----	---	---

	Ind_R	Ind_SAS
RespondentID		
5135740122	50	50
5133300037	15	0
5132253300	40	0
5132096630	35	20
5131990362	60	0

```
In [35]: # descriptive statistics for software preference variables
print('\nDescriptive statistics for survey data -----')
print(software_df.describe())
```

```
Descriptive statistics for survey data -----
```

	My_Java	My_JS	My_Python	My_R	My_SAS	Prof_Java \
count	207.000000	207.000000	207.000000	207.000000	207.000000	207.000000
mean	10.135266	4.797101	31.304348	37.125604	16.637681	9.251208
std	11.383477	6.757764	15.570982	14.576003	13.626400	13.167505
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000	30.000000	5.000000	0.000000
50%	9.000000	0.000000	30.000000	35.000000	15.000000	5.000000
75%	20.000000	10.000000	40.000000	50.000000	25.000000	15.000000
max	70.000000	30.000000	90.000000	100.000000	75.000000	80.000000

	Prof_JS	Prof_Python	Prof_R	Prof_SAS	Ind_Java \
count	207.000000	207.000000	207.000000	207.000000	207.000000
mean	5.840580	30.028986	36.415459	18.463768	11.942029
std	10.812555	19.144802	20.847606	18.831841	14.706399
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	20.000000	25.000000	0.000000	0.000000
50%	0.000000	30.000000	33.000000	15.000000	5.000000
75%	10.000000	40.000000	50.000000	30.000000	20.000000
max	100.000000	100.000000	100.000000	100.000000	70.000000


	Ind_JS	Ind_Python	Ind_R	Ind_SAS
count	207.000000	207.000000	207.000000	207.000000
mean	6.966184	29.772947	32.434783	18.884058
std	10.030721	17.959816	15.912209	19.137623
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	20.000000	22.500000	0.000000
50%	0.000000	30.000000	30.000000	15.000000
75%	10.000000	40.000000	40.000000	30.000000
max	50.000000	95.000000	85.000000	100.000000

```
In [36]: # single scatter plot example
```

```

fig, axis = plt.subplots()
axis.set_xlabel('Personal Preference for R')
axis.set_ylabel('Personal Preference for Python')
plt.title('R and Python Preferences')
scatter_plot = axis.scatter(survey_df['My_R'],
                             survey_df['My_Python'],
                             facecolors = 'none',
                             edgecolors = 'blue')
plt.savefig('plot-scatter-r-python.pdf',
            bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
            orientation='portrait', papertype=None, format=None,
            transparent=True, pad_inches=0.25, frameon=None)

```



```

In [37]: # Creating a list of the survey_df labels which will be used
         # in a for loop to create graphs to compare responses

```

```

survey_df_labels = [
    'Personal Preference for Java/Scala/Spark',
    'Personal Preference for Java/Script/HTML/CSS',
    'Personal Preference for Python',
    'Personal Preference for R',
    'Personal Preference for SAS',
    'Professional Java/Scala/Spark',
    'Professional JavaScript/HTML/CSS',

```

```



'Professional Python',
'Professional R',
'Professional SAS',
'Industry Java/Scala/Spark',
'Industry Java/Script/HTML/CSS',
'Industry Python',
'Industry R',
'Industry SAS'
]



```



```

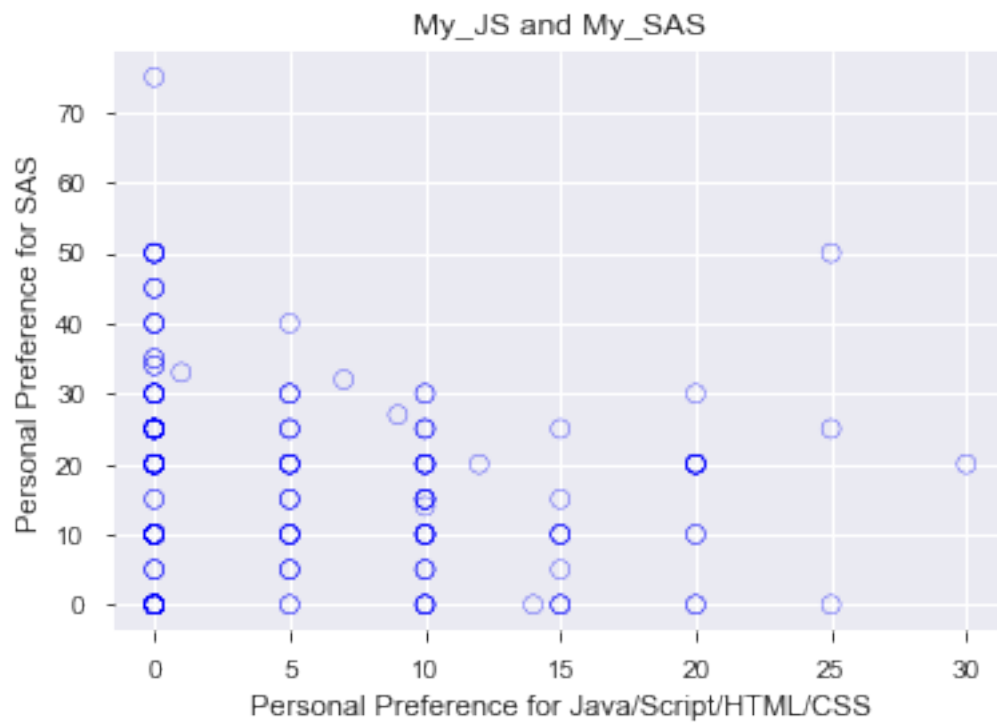

In [38]: # create a set of scatter plots for personal preferences
for i in range(5):
    for j in range(5):
        if i != j:
            file_title = survey_df.columns[i] + '_and_' + survey_df.columns[j]
            plot_title = survey_df.columns[i] + ' and ' + survey_df.columns[j]
            fig, axis = plt.subplots()
            axis.set_xlabel(survey_df_labels[i])
            axis.set_ylabel(survey_df_labels[j])
            plt.title(plot_title)
            scatter_plot = axis.scatter(survey_df[survey_df.columns[i]],
                                         survey_df[survey_df.columns[j]],
                                         facecolors = 'none',
                                         edgecolors = 'blue')
            plt.savefig(file_title + '.pdf',
                        bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                        orientation='portrait', papertype=None, format=None,
                        transparent=True, pad_inches=0.25, frameon=None)



```
































```
In [39]: # examine intercorrelations among software preference variables
# with correlation matrix/heat map
corr_chart(df_corr = software_df)
```

<Figure size 432x288 with 0 Axes>



```
In [40]: corr_matrix = software_df.corr()
corr_matrix["My_Java"].sort_values(ascending=False)
```

```
Out[40]: My_Java      1.000000
Ind_Java      0.728855
Prof_Java     0.697135
My_JS         0.164302
Ind_JS        0.158573
Prof_JS       0.056456
Prof_Python   0.004771
```

```

Ind_Python    -0.115839
My_Python     -0.197282
Prof_SAS      -0.203620
My_SAS        -0.273014
Ind_SAS       -0.278953
Prof_R        -0.290046
Ind_R         -0.307342
My_R          -0.391172
Name: My_Java, dtype: float64

```

0.5 Analysis of New Course Interest and Graduation Date

```

In [41]: # A new dataframe will be created to examine the relationship between
# course interest and expected graduation date.

# Select the graduation date column
example = survey_df.Graduate_Date

# Select the columns for course interest
example_2 = survey_df.iloc[:,15:19]

# Combine the two into a new dataframe and show the first five rows
interest_grad_date_df = pd.concat([example_2, example], axis=1, sort=False)
print(interest_grad_date_df.head())

```

	Python_Course_Interest	Foundations_DE_Course_Interest \
RespondentID		
5135740122	50.0	90.0
5133300037	20.0	50.0
5132253300	100.0	70.0
5132096630	85.0	60.0
5131990362	60.0	10.0

	Analytics_App_Course_Interest	Systems_Analysis_Course_Interest \
RespondentID		
5135740122	51.0	50.0
5133300037	90.0	50.0
5132253300	100.0	60.0
5132096630	90.0	82.0
5131990362	40.0	80.0

	Graduate_Date
RespondentID	
5135740122	NaN
5133300037	Spring 2018
5132253300	Fall 2018
5132096630	Fall 2017
5131990362	Fall 2018

```

In [42]: # Examining the dataframe by grouping by graduation date
         # and checking that this works.

         interest_grad_date_df.groupby(['Graduate_Date']).groups.keys()
         # Out: dict_keys(['2020 or Later', 'Fall 2016', 'Fall 2017', 'Fall 2018',
         # 'Fall 2019', 'Spring 2017', 'Spring 2018', 'Spring 2019', 'Summer 2017', 'Summer 2018',
         # 'Summer 2019', 'Winter 2017', 'Winter 2018', 'Winter 2019'])

Out[42]: dict_keys(['2020 or Later', 'Fall 2016', 'Fall 2017', 'Fall 2018', 'Fall 2019', 'Spring 2017', 'Spring 2018', 'Spring 2019', 'Summer 2017', 'Summer 2018', 'Summer 2019', 'Winter 2017', 'Winter 2018', 'Winter 2019'])

In [43]: # Examining the dataframe by grouping by graduation date
         # and checking that this works.

         len(interest_grad_date_df.groupby(['Graduate_Date']).groups['Fall 2018'])
         #Out: 20

Out[43]: 20

In [44]: # Examine the number of respondents by graduation date

         graduate_date_counts = survey_df.Graduate_Date.value_counts()
         print('\nCount of students by graduation date: ')
         graduate_date_counts

Count of students by graduation date:


Out[44]: Spring 2018      30
         Winter 2017     25
         Winter 2018     25
         Fall 2018       20
         Spring 2017     19
         Fall 2017       14
         Summer 2017     14
         Fall 2016       13
         Summer 2018     11
         Winter 2019     11
         Spring 2019      9
         Fall 2019       5
         2020 or Later    5
         Summer 2019      3
         Name: Graduate_Date, dtype: int64

In [45]: # Examine the number of respondents by graduation date in bar chart

         graduation_date_fig, ax = plt.subplots()

```

```
sns.barplot(y = graduate_date_counts.index,
            x = graduate_date_counts.values, alpha=0.8,
            palette="Blues_d").set_title('Graduation Date of Respondents')
ax.set_xlabel('Number of Occurrences', fontsize=12)
ax.set_ylabel('Quarter', fontsize=12)
graduation_date_fig.savefig('GraduationDate' + '.pdf',
                            bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                            orientation='portrait', papertype=None, format=None,
                            transparent=True, pad_inches=0.25, frameon=None)
```



In [46]: # Examining the dataframe by grouping by graduation date.
The count (number of respondents), min, max, and mean
for each course interest question was evaluated.

```
interest_grad_date_df.groupby(['Graduate_Date']).agg(
    {# find the count min, max, and mean of each course interest column
     'Python_Course_Interest': ['count', min, max, 'mean'],
     'Foundations_DE_Course_Interest': ['count', min, max, 'mean'],
     'Analytics_App_Course_Interest': ['count', min, max, 'mean'],
     'Systems_Analysis_Course_Interest': ['count', min, max, 'mean']})
```

```
Out[46]:
```

	Python_Course_Interest	
	count	min max mean \
Graduate_Date		
2020 or Later	5	35.0 100.0 82.400000

Fall 2016	13	1.0	100.0	76.923077
Fall 2017	14	30.0	100.0	85.000000
Fall 2018	20	15.0	100.0	76.800000
Fall 2019	5	83.0	100.0	91.800000
Spring 2017	18	12.0	100.0	70.166667
Spring 2018	30	0.0	100.0	70.200000
Spring 2019	9	30.0	100.0	77.000000
Summer 2017	14	0.0	100.0	74.642857
Summer 2018	11	0.0	100.0	73.363636
Summer 2019	3	100.0	100.0	100.000000
Winter 2017	25	0.0	100.0	63.520000
Winter 2018	25	5.0	100.0	68.600000
Winter 2019	11	40.0	100.0	73.909091

Foundations_DE_Course_Interest

	count	min	max	mean
Graduate_Date				
2020 or Later	4	25.0	95.0	61.250000
Fall 2016	12	0.0	100.0	70.250000
Fall 2017	14	0.0	100.0	49.142857
Fall 2018	19	10.0	100.0	58.894737
Fall 2019	5	51.0	100.0	80.600000
Spring 2017	18	18.0	100.0	70.444444
Spring 2018	30	0.0	100.0	55.200000
Spring 2019	9	25.0	100.0	62.666667
Summer 2017	14	0.0	100.0	57.500000
Summer 2018	9	0.0	100.0	53.666667
Summer 2019	3	64.0	100.0	88.000000
Winter 2017	25	0.0	100.0	49.920000
Winter 2018	24	0.0	100.0	50.833333
Winter 2019	11	0.0	100.0	56.636364

Analytics_App_Course_Interest

	count	min	max	mean
Graduate_Date				
2020 or Later	5	50.0	90.0	72.200000
Fall 2016	13	0.0	100.0	52.307692
Fall 2017	14	0.0	100.0	60.714286
Fall 2018	19	10.0	100.0	58.210526
Fall 2019	5	20.0	100.0	60.200000
Spring 2017	17	20.0	100.0	59.058824
Spring 2018	30	0.0	100.0	49.133333
Spring 2019	9	30.0	100.0	66.666667
Summer 2017	14	0.0	100.0	59.214286
Summer 2018	10	0.0	100.0	50.500000
Summer 2019	3	100.0	100.0	100.000000
Winter 2017	25	0.0	100.0	43.120000
Winter 2018	25	0.0	100.0	61.800000

Winter 2019	11	0.0	100.0	43.818182
-------------	----	-----	-------	-----------

Systems_Analysis_Course_Interest				
	count	min	max	mean
Graduate_Date				
2020 or Later	5	10.0	70.0	32.000000
Fall 2016	13	3.0	100.0	59.307692
Fall 2017	14	0.0	100.0	61.785714
Fall 2018	18	0.0	100.0	56.611111
Fall 2019	5	10.0	100.0	55.400000
Spring 2017	18	4.0	100.0	55.111111
Spring 2018	29	0.0	100.0	48.689655
Spring 2019	9	30.0	100.0	77.333333
Summer 2017	14	0.0	100.0	52.857143
Summer 2018	9	0.0	100.0	45.555556
Summer 2019	3	90.0	100.0	96.666667
Winter 2017	25	0.0	100.0	40.400000
Winter 2018	25	0.0	100.0	54.280000
Winter 2019	10	0.0	100.0	55.700000

```
In [47]: # Examing the dataframe by grouping by graduation date.
# The mean for each course interest question was evaluated.
# Checking to see if interest changes with graduation date.
```

```
interest_grad_date_df.groupby(['Graduate_Date']).agg(
    [# find the count min, max, and mean of each course interest column
    'Python_Course_Interest': ['mean'],
    'Foundations_DE_Course_Interest': ['mean'],
    'Analytics_App_Course_Interest': ['mean'],
    'Systems_Analysis_Course_Interest': ['mean']})
```


	Python_Course_Interest	Foundations_DE_Course_Interest	\
	mean		mean
Graduate_Date			
2020 or Later	82.400000		61.250000
Fall 2016	76.923077		70.250000
Fall 2017	85.000000		49.142857
Fall 2018	76.800000		58.894737
Fall 2019	91.800000		80.600000
Spring 2017	70.166667		70.444444
Spring 2018	70.200000		55.200000
Spring 2019	77.000000		62.666667
Summer 2017	74.642857		57.500000
Summer 2018	73.363636		53.666667
Summer 2019	100.000000		88.000000
Winter 2017	63.520000		49.920000
Winter 2018	68.600000		50.833333
Winter 2019	73.909091		56.636364

	Analytics_App_Course_Interest	Systems_Analysis_Course_Interest
	mean	mean
Graduate_Date		
2020 or Later	72.200000	32.000000
Fall 2016	52.307692	59.307692
Fall 2017	60.714286	61.785714
Fall 2018	58.210526	56.611111
Fall 2019	60.200000	55.400000
Spring 2017	59.058824	55.111111
Spring 2018	49.133333	48.689655
Spring 2019	66.666667	77.333333
Summer 2017	59.214286	52.857143
Summer 2018	50.500000	45.555556
Summer 2019	100.000000	96.666667
Winter 2017	43.120000	40.400000
Winter 2018	61.800000	54.280000
Winter 2019	43.818182	55.700000

0.6 Analysis of Course Interest

```
In [48]: # Is there any correlation between interest in classes
corr_chart(df_corr = interest_grad_date_df)
```

<Figure size 432x288 with 0 Axes>



In [49]: # Looking at course interest without graduation date using the describe function

```
what_df = survey_df.iloc[:,15:19]
what_df.describe()
```

```
Out[49]:
```


	Python_Course_Interest	Foundations_DE_Course_Interest \
count	206.000000	200.000000
mean	73.529126	58.045000
std	29.835429	32.588079
min	0.000000	0.000000
25%	53.000000	29.500000
50%	82.500000	60.000000
75%	100.000000	89.250000
max	100.000000	100.000000

	Analytics_App_Course_Interest	Systems_Analysis_Course_Interest
count	203.000000	200.000000

mean	55.201970	53.630000
std	34.147954	33.539493
min	0.000000	0.000000
25%	25.000000	21.500000
50%	60.000000	51.500000
75%	85.000000	80.250000
max	100.000000	100.000000

```
In [50]: # Examining the python course interest distribution
python_course_interst_fig, ax = plt.subplots()
sns.distplot(what_df.Python_Course_Interest.dropna()).set_title('Python Course Interest')
python_course_interst_fig.savefig('PythonCourseInterest' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning:
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```



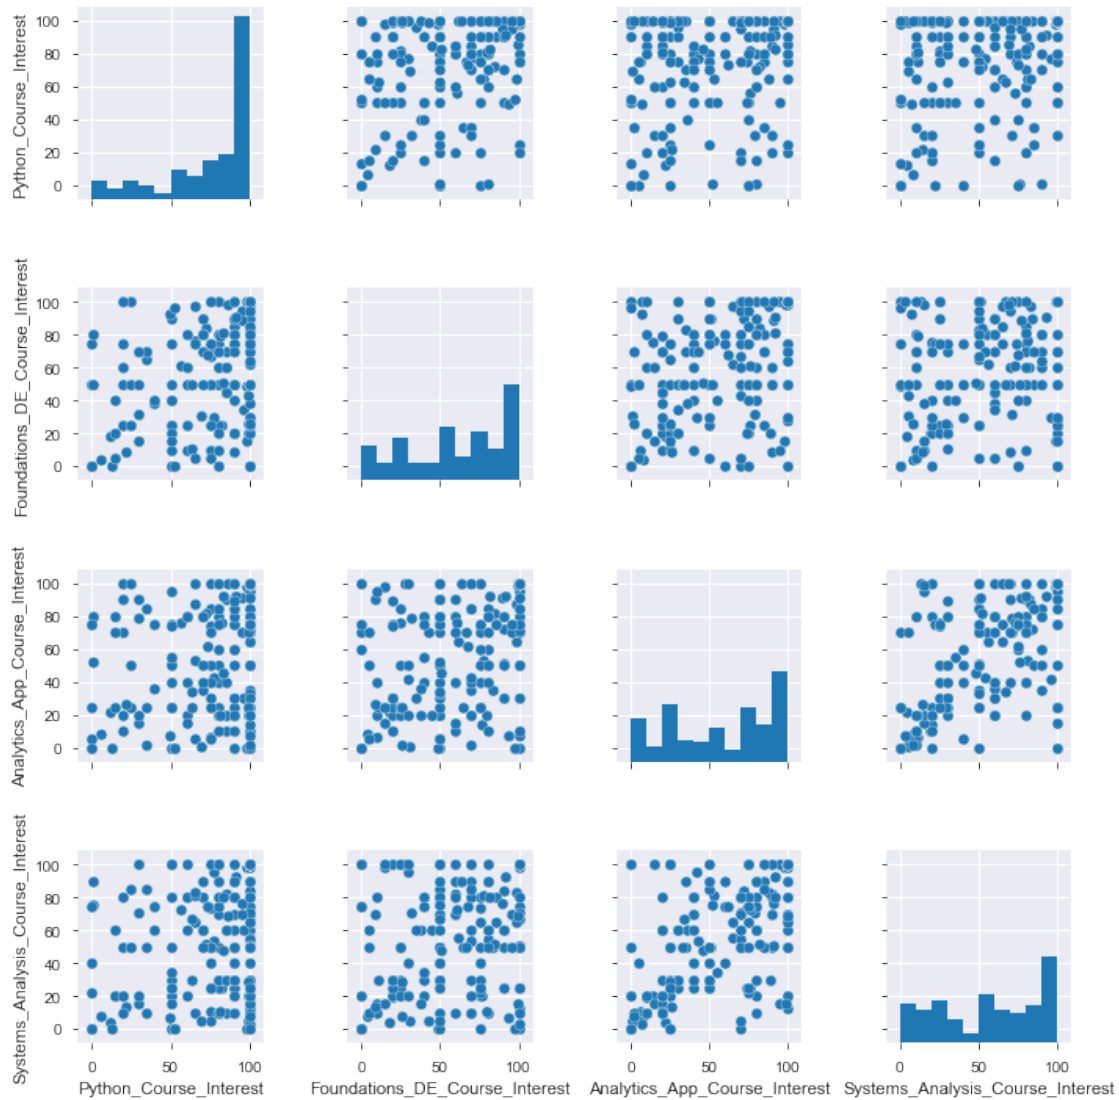
```
In [51]: # Examining the course interest between the four courses via pairplot

course_interest_pairplot_fig = sns.pairplot(what_df.dropna())
course_interest_pairplot_fig.savefig('CourseInterestPairplot' + '.pdf',
```

```

bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
orientation='portrait', papertype=None, format=None,
transparent=True, pad_inches=0.25, frameon=None)

```



0.7 Transformations

```

In [52]: # -----
# transformation code added with version v005
# -----
# transformations a la Scikit Learn
# documentation at http://scikit-learn.org/stable/auto\_examples/preprocessing/plot\_all\_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py
#
#

```

```

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler


# transformations a la Scikit Learn
# select variable to examine, eliminating missing data codes
X = survey_df['Courses_Completed'].dropna()

# Seaborn provides a convenient way to show the effects of transformations
# on the distribution of values being transformed
# Documentation at https://seaborn.pydata.org/generated/seaborn.distplot.html

unscaled_fig, ax = plt.subplots()
sns.distplot(X).set_title('Unscaled')
unscaled_fig.savefig('Transformation-Unscaled' + '.pdf',
                    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                    orientation='portrait', papertype=None, format=None,
                    transparent=True, pad_inches=0.25, frameon=None)

```

/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
 warnings.warn("The 'normed' kwarg is deprecated, and has been "




In [53]: # Reshape your data either using array.reshape(-1, 1) if your data has a single feature
 # or array.reshape(1, -1) if it contains a single sample.

```
X = X.values.reshape(-1, 1)
```


```
standard_fig, ax = plt.subplots()
sns.distplot(StandardScaler().fit_transform(X)).set_title('StandardScaler')
standard_fig.savefig('Transformation-StandardScaler' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
 warnings.warn("The 'normed' kwarg is deprecated, and has been "




```
In [54]: minmax_fig, ax = plt.subplots()
sns.distplot(MinMaxScaler().fit_transform(X)).set_title('MinMaxScaler')
minmax_fig.savefig('Transformation-MinMaxScaler' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
 warnings.warn("The 'normed' kwarg is deprecated, and has been "



```
In [55]: log_fig, ax = plt.subplots()
sns.distplot(np.log(X)).set_title('NaturalLog')
log_fig.savefig('Transformation-NaturalLog' + '.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)
```

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning:
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```




0.8 Transformation of Selected Variable

```
In [56]: # Get rid of NaN values
P = survey_df['Python_Course_Interest'].dropna()
# Change zero to 1 so ln transformation will work
P[P == 0] = 1

# Seaborn provides a convenient way to show the effects of transformations
# on the distribution of values being transformed
# Documentation at https://seaborn.pydata.org/generated/seaborn.distplot.html
```

```
unscaledP_fig, ax = plt.subplots()
sns.distplot(P).set_title('Unscaled')
unscaledP_fig.savefig('Transformation-UnscaledP' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```


```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```



```
In [57]: P = P.values.reshape(-1, 1)
```


```
standardP_fig, ax = plt.subplots()
sns.distplot(StandardScaler().fit_transform(P)).set_title('StandardScaler')
standardP_fig.savefig('Transformation-StandardScalerP' + '.pdf',
    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
    orientation='portrait', papertype=None, format=None,
    transparent=True, pad_inches=0.25, frameon=None)
```

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```

```
In [58]: minmax_figP, ax = plt.subplots()
sns.distplot(MinMaxScaler().fit_transform(P)).set_title('MinMaxScaler')
minmax_figP.savefig('Transformation-MinMaxScalerP' + '.pdf',
                    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                    orientation='portrait', papertype=None, format=None,
                    transparent=True, pad_inches=0.25, frameon=None)
```

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```



```
In [59]: log_figP, ax = plt.subplots()
sns.distplot(np.log(P)).set_title('NaturalLog')
log_figP.savefig('Transformation-NaturalLogP' + '.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)
```

```
/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```

