

Wanat_Assignment_4_final

July 21, 2019

```
In [1]: # Boston Housing Study (Python)
# using data from the Boston Housing Study case
# as described in "Marketing Data Science: Modeling Techniques
# for Predictive Analytics with R and Python" (Miller 2015)

# Here we use data from the Boston Housing Study to evaluate
# regression modeling methods within a cross-validation design.

# program revised by Thomas W. Milller (2017/09/29)

# Scikit Learn documentation for this assignment:
# http://scikit-learn.org/stable/modules/model_evaluation.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.model_selection.KFold.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.linear_model.LinearRegression.html
# http://scikit-learn.org/stable/auto_examples/linear_model/plot_ols.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.linear_model.Ridge.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.linear_model.Lasso.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.linear_model.ElasticNet.html
# http://scikit-learn.org/stable/modules/generated/
#   sklearn.metrics.r2_score.html

# Textbook reference materials:
# Geron, A. 2017. Hands-On Machine Learning with Scikit-Learn
# and TensorFlow. Sebastopol, Calif.: O'Reilly. Chapter 3 Training Models
# has sections covering linear regression, polynomial regression,
# and regularized linear models. Sample code from the book is
# available on GitHub at https://github.com/ageron/handson-ml

# prepare for Python version 3x features and functions
# comment out for Python 3.x execution
# from __future__ import division, print_function
# from future_builtins import ascii, filter, hex, map, oct, zip
```

```

# seed value for random number generators to obtain reproducible results
RANDOM_SEED = 1

# although we standardize X and y variables on input,
# we will fit the intercept term in the models
# Expect fitted values to be close to zero
SET_FIT_INTERCEPT = True

# import base packages into the namespace for this program
import numpy as np
import pandas as pd

# modeling routines from Scikit Learn packages
import sklearn.linear_model
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt # for root mean-squared error calculation
import matplotlib
import matplotlib.pyplot as plt # static plotting
import seaborn as sns # pretty plotting, including heat map
from sklearn.model_selection import train_test_split

```

```

In [2]: # correlation heat map setup for seaborn
def corr_chart(df_corr):
    corr=df_corr.corr()
    #screen top half to get a triangle
    top = np.zeros_like(corr, dtype=np.bool)
    top[np.triu_indices_from(top)] = True
    fig=plt.figure()
    fig, ax = plt.subplots(figsize=(12,12))
    sns.heatmap(corr, mask=top, cmap='coolwarm',
                center = 0, square=True,
                linewidths=.5, cbar_kws={'shrink':.5},
                annot = True, annot_kws={'size': 9}, fmt = '.3f')
    plt.xticks(rotation=45) # rotate variable labels on columns (x axis)
    plt.yticks(rotation=0) # use horizontal variable labels on rows (y axis)
    plt.title('Correlation Heat Map')
    plt.savefig('plot-corr-map.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)

np.set_printoptions(precision=3)

In [3]: # read data for the Boston Housing Study
# creating data frame restdata
boston_input = pd.read_csv('boston.csv')

```

```
# check the pandas DataFrame object boston_input
print('\nboston DataFrame (first and last five rows):')
display(boston_input.head())
display(boston_input.tail())
```

boston DataFrame (first and last five rows):

	neighborhood	crim	zn	indus	chas	nox	rooms	age	dis	rad	\
0	Nahant	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	
1	Swampscott	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	
2	Swampscott	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	
3	Marblehead	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	
4	Marblehead	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	

	tax	ptratio	lstat	mv
0	296	15.3	4.98	24.0
1	242	17.8	9.14	21.6
2	242	17.8	4.03	34.7
3	222	18.7	2.94	33.4
4	222	18.7	5.33	36.2

	neighborhood	crim	zn	indus	chas	nox	rooms	age	dis	rad	\
501	Winthrop	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	
502	Winthrop	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	
503	Winthrop	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	
504	Winthrop	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	
505	Winthrop	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	

	tax	ptratio	lstat	mv
501	273	21.0	9.67	22.4
502	273	21.0	9.08	20.6
503	273	21.0	5.64	23.9
504	273	21.0	6.48	22.0
505	273	21.0	7.88	19.0

```
In [4]: print('\nGeneral description of the boston_input DataFrame:')
print(boston_input.info())
```

```
General description of the boston_input DataFrame:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
neighborhood    506 non-null object
```

```

crim          506 non-null float64
zn            506 non-null float64
indus         506 non-null float64
chas          506 non-null int64
nox           506 non-null float64
rooms         506 non-null float64
age           506 non-null float64
dis           506 non-null float64
rad           506 non-null int64
tax           506 non-null int64
ptratio       506 non-null float64
lstat         506 non-null float64
mv            506 non-null float64
dtypes: float64(10), int64(3), object(1)
memory usage: 55.4+ KB
None

```

```

In [5]: # drop neighborhood from the data being considered
        boston = boston_input.drop('neighborhood', 1)
        print('\nGeneral description of the boston DataFrame:')
        print(boston.info())

```

```

General description of the boston DataFrame:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 13 columns):
crim          506 non-null float64
zn            506 non-null float64
indus         506 non-null float64
chas          506 non-null int64
nox           506 non-null float64
rooms         506 non-null float64
age           506 non-null float64
dis           506 non-null float64
rad           506 non-null int64
tax           506 non-null int64
ptratio       506 non-null float64
lstat         506 non-null float64
mv            506 non-null float64
dtypes: float64(10), int64(3)
memory usage: 51.5 KB
None

```

```

In [6]: print('\nDescriptive statistics of the boston DataFrame:')
        print(boston.describe())

```

Descriptive statistics of the boston DataFrame:


	crim	zn	indus	chas	nox	rooms \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677082	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	age	dis	rad	tax	ptratio	lstat \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	12.653063
std	28.148861	2.105710	8.707259	168.537116	2.164946	7.141062
min	2.900000	1.129600	1.000000	187.000000	12.600000	1.730000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	6.950000
50%	77.500000	3.207450	5.000000	330.000000	19.050000	11.360000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	16.955000
max	100.000000	12.126500	24.000000	711.000000	22.000000	37.970000

	mv
count	506.000000
mean	22.528854
std	9.182176
min	5.000000
25%	17.025000
50%	21.200000
75%	25.000000
max	50.000000

In [7]: corr_chart(boston)

<Figure size 432x288 with 0 Axes>



```
In [8]: # set up preliminary data for data for fitting the models
# the first column is the median housing value response
# the remaining columns are the explanatory variables
prelim_model_data = np.array([boston.mv,\
    boston.crim,\
    boston.zn,\
    boston.indus,\
    boston.chas,\
    boston.nox,\
    boston.rooms,\
    boston.age,\
    boston.dis,\
    boston.rad,\
    boston.tax,\
    boston.ptratio,\
    boston.lstat]).T
```

```
In [9]: # dimensions of the polynomial model X input and y response
        # preliminary data before standardization
        print('\nData dimensions:', prelim_model_data.shape)
```

Data dimensions: (506, 13)

```
In [10]: # standard scores for the columns... along axis 0
         from sklearn.preprocessing import StandardScaler
         scaler = StandardScaler()
         print(scaler.fit(prelim_model_data))
```

StandardScaler(copy=True, with_mean=True, with_std=True)

```
In [11]: # show standardization constants being employed
         print(scaler.mean_)
         print(scaler.scale_)
```

```
[2.253e+01 3.614e+00 1.136e+01 1.114e+01 6.917e-02 5.547e-01 6.285e+00
 6.857e+01 3.795e+00 9.549e+00 4.082e+02 1.846e+01 1.265e+01]
[9.173e+00 8.593e+00 2.330e+01 6.854e+00 2.537e-01 1.158e-01 7.019e-01
 2.812e+01 2.104e+00 8.699e+00 1.684e+02 2.163e+00 7.134e+00]
```

```
In [12]: # the model data will be standardized form of preliminary model data
         model_data = scaler.fit_transform(prelim_model_data)


         # dimensions of the polynomial model X input and y response
         # all in standardized units of measure
         print('\nDimensions for model_data:', model_data.shape)
```

Dimensions for model_data: (506, 13)

```
In [13]: sns.distplot(boston.mv)
```

/Users/jmwanat/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning
warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x110f61a90>



In [14]: *#split data and response*

```
boston_X = boston.drop('mv', axis=1)
boston_y = boston.mv.copy()
```

In [15]: boston_X.head()

```
Out[15]:
```

	crim	zn	indus	chas	nox	rooms	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	


```
lstat
0    4.98
1    9.14
2    4.03
3    2.94
4    5.33
```

In [16]: boston_y.head()

```
Out[16]: 0    24.0
         1    21.6
```



```

2    34.7
3    33.4
4    36.2
Name: mv, dtype: float64

```

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(boston_X, boston_y, random_state=)
```

0.1 Linear Regression

```
In [18]: from sklearn.linear_model import LinearRegression
```

```

lin_reg = LinearRegression()
lr = lin_reg.fit(X_train, y_train)
lr_intercept = lr.intercept_
lr_coef = lr.coef_
print('\n-----')
print('Linear Regression\n')
print('intercept:', lr_intercept)
print('coefficients:', lr_coef)

```

```

-----
Linear Regression

```

```

intercept: 48.18770627941787
coefficients: [-1.195e-01  5.937e-02  3.721e-02  2.529e+00 -2.183e+01  2.780e+00
 7.906e-03 -1.524e+00  2.907e-01 -1.136e-02 -9.388e-01 -5.818e-01]

```

```
In [19]: print("Training set score: {:.2f}".format(lr.score(X_train, y_train)))
        print("Test set score: {:.2f}".format(lr.score(X_test, y_test)))
```

```

Training set score: 0.72
Test set score: 0.77

```

```
In [20]: from sklearn.model_selection import cross_val_score
        scores = cross_val_score(lin_reg, boston_X, boston_y,
                                scoring="neg_mean_squared_error", cv=10)
        lin_rmse_scores = np.sqrt(-scores)
```

```
In [21]: def display_scores(scores):
        print("Scores:", scores)
        print("Mean:", scores.mean())
        print("Standard deviation:", scores.std())

        print('RMSE Linear Regression')
        display_scores(lin_rmse_scores)

```

RMSE Linear Regression


Scores: [2.826 3.806 4.026 5.983 5.704 4.55 3.154 12.3 6.143 3.056]

Mean: 5.154728385946134

Standard deviation: 2.6514954341431514

```
In [22]: # predicting the test set results
y_pred_lr = lin_reg.predict(X_test)
```

```
In [23]: # Plotting Scatter graph to show the prediction
# results - 'ytrue' value vs 'y_pred' value
plt.scatter(y_test, y_pred_lr, c = 'blue')
plt.xlabel("Price: in $1000's")
plt.ylabel("Predicted value")
plt.title("True value vs predicted value : Linear Regression")
plt.savefig('true_vs_predicted_LR.pdf')
plt.show()
```



0.2 Ridge Regression

```
In [24]: from sklearn.linear_model import Ridge
```

```
# Linear least squares with l2 regularization.
# alpha = Regularization strength; must be a positive float.
```

```
# Regularization improves the conditioning of the problem
# and reduces the variance of the estimates.
# Larger values of alpha specify stronger regularization.
```

```
ridge = Ridge()
rr = ridge.fit(X_train, y_train)
rr_intercept = ridge.intercept_
rr_coef = ridge.coef_

print('\n-----')
print('Ridge Regression\n')
print('intercept:', rr_intercept)
print('coefficients:', rr_coef)
```

```
-----
Ridge Regression
```

```
intercept: 41.23001324745856
coefficients: [-1.149e-01  6.041e-02 -9.447e-03  2.277e+00 -1.182e+01  2.869e+00
 -3.949e-04 -1.386e+00  2.698e-01 -1.250e-02 -8.179e-01 -5.947e-01]
```

```
In [25]: print('Ridge alpha = 1')
         print("Training set score: {:.2f}".format(ridge.score(X_train, y_train)))
         print("Test set score: {:.2f}".format(ridge.score(X_test, y_test)))
```

```
Ridge alpha = 1
Training set score: 0.72
Test set score: 0.77
```

```
In [26]: ridge_scores = cross_val_score(ridge, boston_X, boston_y,
                                         scoring="neg_mean_squared_error", cv=10)
         ridge01_rmse_scores = np.sqrt(-ridge_scores)
         print('RMSE Ridge Regression alpha = 1')
         display_scores(ridge01_rmse_scores)
```

```
RMSE Ridge Regression alpha = 1
Scores: [ 2.846  3.584  3.529  6.117  5.509  4.401  3.081 12.268  6.305  3.207]
Mean: 5.0848231944069475
Standard deviation: 2.683888647423676
```

```
In [27]: ridge10 = Ridge(alpha=10).fit(X_train, y_train)
         print('Ridge alpha = 10')
         print("Training set score: {:.2f}".format(ridge10.score(X_train, y_train)))
         print("Test set score: {:.2f}".format(ridge10.score(X_test, y_test)))
```

```
Ridge alpha = 10
Training set score: 0.70
Test set score: 0.77
```

```
In [28]: ridge01 = Ridge(alpha=0.1).fit(X_train, y_train)
         print('Ridge alpha = 0.1')
         print("Training set score: {:.2f}".format(ridge01.score(X_train, y_train)))
         print("Test set score: {:.2f}".format(ridge01.score(X_test, y_test)))
```


```
Ridge alpha = 0.1
Training set score: 0.72
Test set score: 0.77
```

```
In [29]: ridge_scores = cross_val_score(ridge01, boston_X, boston_y,
                                         scoring="neg_mean_squared_error", cv=10)
         ridge01_rmse_scores = np.sqrt(-ridge_scores)
         print('RMSE Ridge Regression alpha = 0.01')
         display_scores(ridge01_rmse_scores)
```

```
RMSE Ridge Regression alpha = 0.01
Scores: [ 2.826  3.768  3.919  6.002  5.67   4.523  3.135 12.297  6.169  3.074]
Mean: 5.138375588249978
Standard deviation: 2.658992425555945
```

```
In [30]: # predicting the test set results
         y_pred_rr = ridge.predict(X_test)
```

```
In [31]: # Plotting Scatter graph to show the prediction
         # results - 'ytrue' value vs 'y_pred' value
         plt.scatter(y_test, y_pred_rr, c = 'red')
         plt.xlabel("Price: in $1000's")
         plt.ylabel("Predicted value")
         plt.title("True value vs predicted value : Ridge Regression")
         plt.savefig('true_vs_predicted_RR.pdf')
         plt.show()
```



0.3 Lasso

In [32]: `from sklearn.linear_model import Lasso`

```
# Linear Model trained with L1 prior as regularizer
# alpha = Regularization strength, constant that multiplies the L1 term.
```

```
lasso = Lasso()
lasso.fit(X_train, y_train)
lasso.intercept_, lasso.coef_
```

```
lasso_intercept = lasso.intercept_
lasso_coef = lasso.coef_
```

```
print('\n-----')
print('Lasso with default alpha = 1.0\n')
print('intercept:', lasso_intercept)
print('coefficients:', lasso_coef)
```

```
-----
Lasso with default alpha = 1.0
```

```
intercept: 48.83289146303929
```

```
coefficients: [-0.073  0.057 -0.      0.    -0.      0.      0.023 -0.587  0.23  -0.014
 -0.67  -0.812]
```

```
In [33]: print('Lasso alpha = 1, max_iter = 1000')
        print("Training set score: {:.2f}".format(lasso.score(X_train, y_train)))
        print("Test set score: {:.2f}".format(lasso.score(X_test, y_test)))
        print("Number of features used:", np.sum(lasso.coef_ != 0))
```

```
Lasso alpha = 1, max_iter = 1000
Training set score: 0.65
Test set score: 0.66
Number of features used: 8
```

```
In [34]: lasso_scores = cross_val_score(lasso, boston_X, boston_y,
                                         scoring="neg_mean_squared_error", cv=10)
        lasso_rmse_scores = np.sqrt(-lasso_scores)
        print('RMSE Lasso alpha = 1, max_iter = 1000')
        display_scores(lasso_rmse_scores)
```

```
RMSE Lasso alpha = 1, max_iter = 1000
Scores: [ 3.359  4.247  3.411  7.813  6.856  6.478  4.231 10.109  5.276  3.452]
Mean: 5.523163553242023
Standard deviation: 2.1390100357462694
```

```
In [35]: # we increase the default setting of "max_iter",
         # otherwise the model would warn us that we should increase max_iter.
         # regularization parameter, alpha, that controls
         # how strongly coefficients are pushed toward zero.
         # To reduce underfitting, lets try decreasing alpha.

        lasso001 = Lasso(alpha=0.01, max_iter=100000).fit(X_train, y_train)
        print('Lasso alpha = 0.01, max_iter = 100000')
        print("Training set score: {:.2f}".format(lasso001.score(X_train, y_train)))
        print("Test set score: {:.2f}".format(lasso001.score(X_test, y_test)))
        print("Number of features used:", np.sum(lasso001.coef_ != 0))
```


```
Lasso alpha = 0.01, max_iter = 100000
Training set score: 0.72
Test set score: 0.77
Number of features used: 12
```

```
In [36]: lasso001_scores = cross_val_score(lasso001, boston_X, boston_y,
                                         scoring="neg_mean_squared_error", cv=10)
        lasso001_rmse_scores = np.sqrt(-lasso001_scores)
        print('RMSE Lasso alpha = 0.01, max_iter = 100000')
        display_scores(lasso001_rmse_scores)
```

```
RMSE Lasso alpha = 0.01, max_iter = 100000
Scores: [ 2.827  3.729  3.789  6.051  5.606  4.492  3.114 12.306  6.231  3.114]
Mean: 5.1259392630655345
Standard deviation: 2.671885020579878
```

```
In [37]: # predicting the test set results
        y_pred_lasso = lasso001.predict(X_test)

In [38]: # Plotting Scatter graph to show the prediction
        # results - 'ytrue' value vs 'y_pred' value
        plt.scatter(y_test, y_pred_lasso, c = 'purple')
        plt.xlabel("Price: in $1000's")
        plt.ylabel("Predicted value")
        plt.title("True value vs predicted value : Lasso")
        plt.savefig('true_vs_predicted_lasso.pdf')
        plt.show()
```



0.4 Elastic Net

```
In [39]: from sklearn.linear_model import ElasticNet

        # Linear regression with combined L1 and L2 priors as regularizer
        # alpha = Regularization strength, constant that multiplies the penalty terms.
```

```

elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5)
elastic_net.fit(X_train, y_train)

elastic_net_intercept = elastic_net.intercept_
elastic_net_coef = elastic_net.coef_

print('\n-----')
print('Elastic Net\n')
print('intercept:', elastic_net_intercept)
print('coefficients:', elastic_net_coef)

-----
Elastic Net

intercept: 37.13744686978241
coefficients: [-0.11  0.063 -0.055  0.912 -0.307  2.443 -0.004 -1.179  0.264 -0.015
 -0.709 -0.649]

In [40]: print('alpha=1.0, l1_ratio=0.5')
         print("Training set score: {:.2f}".format(elastic_net.score(X_train, y_train)))
         print("Test set score: {:.2f}".format(elastic_net.score(X_test, y_test)))

alpha=1.0, l1_ratio=0.5
Training set score: 0.70
Test set score: 0.76

In [41]: elastic_net_scores = cross_val_score(elastic_net, boston_X, boston_y,
                                              scoring="neg_mean_squared_error", cv=10)
         elastic_net_rmse_scores = np.sqrt(-elastic_net_scores)
         print('RMSE Elastic Net alpha=1.0, l1_ratio=0.5')
         display_scores(elastic_net_rmse_scores)

RMSE Elastic Net alpha=1.0, l1_ratio=0.5
Scores: [ 2.985  3.518  3.336  6.524  5.516  4.506  3.313 11.316  6.237  3.554]
Mean: 5.080429131773322
Standard deviation: 2.4115057536950926

In [42]: elastic_net_2 = ElasticNet(alpha=0.5, l1_ratio=0.1)
         elastic_net_2.fit(X_train, y_train)
         print('alpha=0.5, l1_ratio=0.1')
         print("Training set score: {:.2f}".format(elastic_net.score(X_train, y_train)))
         print("Test set score: {:.2f}".format(elastic_net.score(X_test, y_test)))

```



```
alpha=0.5, l1_ratio=0.1
Training set score: 0.70
Test set score: 0.76
```

```
In [43]: elastic_net_2_scores = cross_val_score(elastic_net_2, boston_X, boston_y,
                                                scoring="neg_mean_squared_error", cv=10)
        elastic_net_2_rmse_scores = np.sqrt(-elastic_net_2_scores)
        print('RMSE Elastic Net alpha=0.5, l1_ratio=0.1')
        display_scores(elastic_net_2_rmse_scores)
```

```
RMSE Elastic Net alpha=0.5, l1_ratio=0.1
Scores: [ 2.985  3.518  3.336  6.524  5.516  4.506  3.313 11.316  6.237  3.554]
Mean: 5.080429131773322
Standard deviation: 2.4115057536950926
```

```
In [44]: # predicting the test set results
        y_pred_en = elastic_net.predict(X_test)
```

```
In [45]: # Plotting Scatter graph to show the prediction
        # results - 'ytrue' value vs 'y_pred' value
        plt.scatter(y_test, y_pred_en, c = 'orange')
        plt.xlabel("Price: in $1000's")
        plt.ylabel("Predicted value")
        plt.title("True value vs predicted value : Elastic Net")
        plt.savefig('true_vs_predicted_EN.pdf')
        plt.show()
```



0.5 Grid Search

0.5.1 Use Grid Search to find optimal parameters for Elastic Net

```
In [46]: from sklearn.model_selection import GridSearchCV
```

```
# GridSearchCV: specify which hyperparameters you want  
# it to experiment with, and what values to try out,  
# and it will evaluate all the possible combinations  
# of hyperparameter values, using cross-validation
```

```
param_grid = [  
    {'alpha': [0.1, 0.5, 1, 10, 20], 'l1_ratio': [0.1, 0.25, 0.5, 0.75]},  
    ]
```

```
elastic_net_gs = ElasticNet()
```

```
grid_search = GridSearchCV(elastic_net_gs, param_grid, cv=5,  
                           scoring='neg_mean_squared_error',  
                           return_train_score=True)
```

```
grid_search.fit(boston_X, boston_y)
```

```
Out[46]: GridSearchCV(cv=5, error_score='raise',  
    estimator=ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True, l1_ratio=0.5,  
    max_iter=1000, normalize=False, positive=False, precompute=False,  
    random_state=None, selection='cyclic', tol=0.0001, warm_start=False),  
    fit_params=None, iid=True, n_jobs=1,  
    param_grid=[{'alpha': [0.1, 0.5, 1, 10, 20], 'l1_ratio': [0.1, 0.25, 0.5, 0.75]},  
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,  
    scoring='neg_mean_squared_error', verbose=0)
```

```
In [47]: print('\n-----')  
print('The parameters chosen by Grid Search for Elastic Net:\n')  
grid_search.best_params_
```

```
-----  
The parameters chosen by Grid Search for Elastic Net:
```

```
Out[47]: {'alpha': 0.5, 'l1_ratio': 0.1}
```

```
In [48]: grid_search.best_estimator_
```

```
Out[48]: ElasticNet(alpha=0.5, copy_X=True, fit_intercept=True, l1_ratio=0.1,
                    max_iter=1000, normalize=False, positive=False, precompute=False,
                    random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
```

```
In [49]: cvres = grid_search.cv_results_
```

```
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
    print(np.sqrt(-mean_score), params)
```

```
5.6464930839855 {'alpha': 0.1, 'l1_ratio': 0.1}
5.677263610425171 {'alpha': 0.1, 'l1_ratio': 0.25}
5.740279703029267 {'alpha': 0.1, 'l1_ratio': 0.5}
5.818655787231755 {'alpha': 0.1, 'l1_ratio': 0.75}
5.546062243897276 {'alpha': 0.5, 'l1_ratio': 0.1}
5.5740235714285875 {'alpha': 0.5, 'l1_ratio': 0.25}
5.630446396462518 {'alpha': 0.5, 'l1_ratio': 0.5}
5.727868299091872 {'alpha': 0.5, 'l1_ratio': 0.75}
5.613846498566613 {'alpha': 1, 'l1_ratio': 0.1}
5.649877635818467 {'alpha': 1, 'l1_ratio': 0.25}
5.729059741699742 {'alpha': 1, 'l1_ratio': 0.5}
5.832027373152929 {'alpha': 1, 'l1_ratio': 0.75}
6.384233839729841 {'alpha': 10, 'l1_ratio': 0.1}
6.568917497795396 {'alpha': 10, 'l1_ratio': 0.25}
6.684133591183442 {'alpha': 10, 'l1_ratio': 0.5}
6.7331033985792015 {'alpha': 10, 'l1_ratio': 0.75}
6.794337009151011 {'alpha': 20, 'l1_ratio': 0.1}
6.900179452400165 {'alpha': 20, 'l1_ratio': 0.25}
7.0171707864846065 {'alpha': 20, 'l1_ratio': 0.5}
7.190116716832263 {'alpha': 20, 'l1_ratio': 0.75}
```

```
In [50]: elastic_net_try = ElasticNet(alpha=0.5, l1_ratio=0.1, max_iter=1000)
elastic_net_try.fit(X_train, y_train)
```

```
Out[50]: ElasticNet(alpha=0.5, copy_X=True, fit_intercept=True, l1_ratio=0.1,
                    max_iter=1000, normalize=False, positive=False, precompute=False,
                    random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
```

```
In [51]: elastic_net_try.fit(X_train, y_train)
print('Elastic Net alpha=0.5, l1_ratio=0.1')
print("Training set score: {:.2f}".format(elastic_net_try.score(X_train, y_train)))
print("Test set score: {:.2f}".format(elastic_net_try.score(X_test, y_test)))
```

```
Elastic Net alpha=0.5, l1_ratio=0.1
Training set score: 0.68
Test set score: 0.71
```

```
In [52]: elastic_net_try_scores = cross_val_score(elastic_net_try, boston_X, boston_y,
                                                  scoring="neg_mean_squared_error", cv=10)
```

```

elastic_net_try_rmse_scores = np.sqrt(-elastic_net_try_scores)
print('RMSE Elastic Net alpha=0.5, l1_ratio=0.1')
display_scores(elastic_net_2_rmse_scores)

```

RMSE Elastic Net alpha=0.5, l1_ratio=0.1

Scores: [2.985 3.518 3.336 6.524 5.516 4.506 3.313 11.316 6.237 3.554]

Mean: 5.080429131773322

Standard deviation: 2.4115057536950926

```


In [53]: # predicting the test set results
y_pred = elastic_net_try.predict(X_test)

```

```

In [54]: # Plotting Scatter graph to show the prediction
# results - 'ytrue' value vs 'y_pred' value
plt.scatter(y_test, y_pred, c = 'green')
plt.xlabel("Price: in $1000's")
plt.ylabel("Predicted value")
plt.title("True value vs predicted value : Elastic Net Optimized")
plt.savefig('true_vs_predicted_EN_optimized.pdf')
plt.show()

```



0.6 Pipeline

0.6.1 Use pipeline to evaluate Random Forest and Gradient Boosting Regressors

```
In [55]: from sklearn.pipeline import Pipeline
         from sklearn.ensemble import RandomForestRegressor
         from sklearn.ensemble import GradientBoostingRegressor

In [56]: #For Gradient Boosting Regressor:
         #The learning_rate hyperparameter scales the contribution of each tree.
         #If you set it to a low value, such as 0.1, you will need more trees in
         #the ensemble to fit the training set, but the predictions will usually
         #generalize better. This is a regularization technique called shrinkage.

         # create a list of classifiers to be used in the pipeline
en_classifiers = [
    RandomForestRegressor(n_estimators=100, max_leaf_nodes=16, n_jobs=-1, bootstrap=True),
    GradientBoostingRegressor(n_estimators=100, max_depth=2, learning_rate=1.0)]

In [57]: #loop through each of the classifiers in the list
         # use pipeline to instantiate the model and fit to training data
         # score each model on training and test data sets
         # use cross validation score to calculate RMSE

for classifier in en_classifiers:
    pipe = Pipeline(steps=[('classifier', classifier)])
    pipe.fit(X_train, y_train)
    pipe_scores = cross_val_score(pipe, boston_X, boston_y,
                                   scoring="neg_mean_squared_error", cv=10)
    pipe_rmse_scores = np.sqrt(-pipe_scores)
    print('\n-----')
    print('\nClassifier evaluation for:')
    print(classifier)
    print('Model training score: {:.3f}'.format(pipe.score(X_train, y_train)))
    print('Model test score: {:.3f}'.format(pipe.score(X_test, y_test)))
    print('RMSE mean: {:.3f}'.format(pipe_rmse_scores.mean()))
    print('RMSE standard deviation: {:.3f}'.format(pipe_rmse_scores.std()))

-----

Classifier evaluation for:
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                      max_features='auto', max_leaf_nodes=16,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
                      oob_score=False, random_state=1, verbose=0, warm_start=False)
Model training score: 0.932
```

Model test score: 0.895
RMSE mean: 4.269
RMSE standard deviation: 2.058

Classifier evaluation for:

```
GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,  
    learning_rate=1.0, loss='ls', max_depth=2, max_features=None,  
    max_leaf_nodes=None, min_impurity_decrease=0.0,  
    min_impurity_split=None, min_samples_leaf=1,  
    min_samples_split=2, min_weight_fraction_leaf=0.0,  
    n_estimators=100, presort='auto', random_state=None,  
    subsample=1.0, verbose=0, warm_start=False)
```

Model training score: 0.996
Model test score: 0.801
RMSE mean: 5.025
RMSE standard deviation: 1.566

0.7 Grid Search

0.7.1 Grid Search for Random Forest Regressor

```
In [58]: # GridSearchCV: specify which parameters you want  
        # it to experiment with, and what values to try out,  
        # and it will evaluate all the possible combinations  
        # of parameter values, using cross-validation  
        # to find the optimal settings  
  
        # Information on parameters:  
        #min_samples_split (the minimum number of samples a node must have before it can be split),  
        #min_samples_leaf (the minimum number of samples a leaf node must have),  
        #min_weight_fraction_leaf (same as min_samples_leaf but expressed as a  
        #                          fraction of the total number of weighted instances),  
        #max_leaf_nodes (maximum number of leaf nodes),  
        #max_features (maximum number of features that are  
        #               evaluated for splitting at each node).  
        #Increasing min_* hyperparameters or reducing max_* hyperparameters will regularize the model.  
        #Muller and Guido suggest using max_features='log2' for random forest regression problems  
        #and max_features='sqrt' for classification problems.  
  
In [59]: param_grid = [  
        {'max_leaf_nodes': [2,6,10,14,16,20],  
         'max_features': ['auto', 'log2']}]
```

```
rfr = RandomForestRegressor(n_jobs=-1, bootstrap=True, random_state=RANDOM_SEED)
```

```
rfr_grid_search = GridSearchCV(rfr, param_grid, cv=5,
                               scoring='neg_mean_squared_error',
                               return_train_score=True)
```

```
rfr_grid_search.fit(boston_X, boston_y)
```

```
Out [59]: GridSearchCV(cv=5, error_score='raise',
                      estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=-1,
oob_score=False, random_state=1, verbose=0, warm_start=False),
fit_params=None, iid=True, n_jobs=1,
param_grid=[{'max_leaf_nodes': [2, 6, 10, 14, 16, 20], 'max_features': ['auto']},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring='neg_mean_squared_error', verbose=0)
```

```
In [60]: print('\n-----')
print('The parameters chosen by Grid Search for Random Forest Regressor:\n ')
rfr_grid_search.best_params_
```

The parameters chosen by Grid Search for Random Forest Regressor:

```
Out [60]: {'max_features': 'auto', 'max_leaf_nodes': 20}
```

```
In [61]: #print the mean RMSE for each model evaluated
#the model with the lowest mean RMSE is best
rfr_cvres = rfr_grid_search.cv_results_

for mean_score, params in zip(rfr_cvres["mean_test_score"], rfr_cvres["params"]):
    print(np.sqrt(-mean_score), params)
```

```
7.3111621613506115 {'max_features': 'auto', 'max_leaf_nodes': 2}
5.357874860005873 {'max_features': 'auto', 'max_leaf_nodes': 6}
5.149457664123636 {'max_features': 'auto', 'max_leaf_nodes': 10}
5.02827504240986 {'max_features': 'auto', 'max_leaf_nodes': 14}
4.975451204037436 {'max_features': 'auto', 'max_leaf_nodes': 16}
4.943763789163307 {'max_features': 'auto', 'max_leaf_nodes': 20}
7.972419380044822 {'max_features': 'log2', 'max_leaf_nodes': 2}
5.540629257524592 {'max_features': 'log2', 'max_leaf_nodes': 6}
5.186832486196532 {'max_features': 'log2', 'max_leaf_nodes': 10}
5.093682534659262 {'max_features': 'log2', 'max_leaf_nodes': 14}
5.035181888605574 {'max_features': 'log2', 'max_leaf_nodes': 16}
```

```
4.989562939487351 {'max_features': 'log2', 'max_leaf_nodes': 20}
```

```
In [62]: #perform grid search again to find optimal max_depth
```

```
param_grid = [  
    {'max_leaf_nodes': [20],  
     'max_features': ['auto'],  
     'max_depth': [2,6,8,10,12]}]  
  
rfr = RandomForestRegressor(n_jobs=-1, bootstrap=True, random_state=RANDOM_SEED)  
  
rfr_grid_search = GridSearchCV(rfr, param_grid, cv=5,  
                               scoring='neg_mean_squared_error',  
                               return_train_score=True)  
  
rfr_grid_search.fit(boston_X, boston_y)
```

```
Out[62]: GridSearchCV(cv=5, error_score='raise',  
                      estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,  
                                                       max_features='auto', max_leaf_nodes=None,  
                                                       min_impurity_decrease=0.0, min_impurity_split=None,  
                                                       min_samples_leaf=1, min_samples_split=2,  
                                                       min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=-1,  
                                                       oob_score=False, random_state=1, verbose=0, warm_start=False),  
                      fit_params=None, iid=True, n_jobs=1,  
                      param_grid=[{'max_leaf_nodes': [20], 'max_features': ['auto'], 'max_depth': [2,  
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,  
scoring='neg_mean_squared_error', verbose=0)
```

```
In [63]: print('\n-----')  
         print('The parameters choosen by Grid Search for Random Forest Regressor:\n ')  
         rfr_grid_search.best_params_
```

```
-----  
The parameters choosen by Grid Search for Random Forest Regressor:
```

```
Out[63]: {'max_depth': 6, 'max_features': 'auto', 'max_leaf_nodes': 20}
```

0.7.2 Grid Search for Gradient Boosting Regressor

```
In [64]: gbr_param_grid = [  
    {'max_leaf_nodes': [2,6,10,14,16,20],  
     'max_features': ['auto', 'log2']}]  
  
#max_depth=2,
```



```
gbr = GradientBoostingRegressor(n_estimators=100, learning_rate=1.0)
```

```
gbr_grid_search = GridSearchCV(gbr, gbr_param_grid, cv=5,  
                                scoring='neg_mean_squared_error',  
                                return_train_score=True)
```

```
gbr_grid_search.fit(boston_X, boston_y)
```

```
Out [64]: GridSearchCV(cv=5, error_score='raise',  
                      estimator=GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=  
                        learning_rate=1.0, loss='ls', max_depth=3, max_features=None,  
                        max_leaf_nodes=None, min_impurity_decrease=0.0,  
                        min_impurity_split=None, min_samples_leaf=1,  
                        min_samples_split=2, min_weight_fraction_leaf=0.0,  
                        n_estimators=100, presort='auto', random_state=None,  
                        subsample=1.0, verbose=0, warm_start=False),  
                      fit_params=None, iid=True, n_jobs=1,  
                      param_grid=[{'max_leaf_nodes': [2, 6, 10, 14, 16, 20], 'max_features': ['auto']},  
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,  
                      scoring='neg_mean_squared_error', verbose=0)
```

```
In [65]: print('\n-----')  
         print('The parameters chosen by Grid Search for Random Forest Regressor:\n ')  
         gbr_grid_search.best_params_
```

```
-----  
The parameters chosen by Grid Search for Random Forest Regressor:
```

```
Out [65]: {'max_features': 'log2', 'max_leaf_nodes': 2}
```

```
In [66]: #print the mean RMSE for each model evaluated  
         #the model with the lowest mean RMSE is best
```

```
gbr_cvres = gbr_grid_search.cv_results_
```

```
for mean_score, params in zip(gbr_cvres["mean_test_score"], gbr_cvres["params"]):  
    print(np.sqrt(-mean_score), params)
```

```
5.72552465589504 {'max_features': 'auto', 'max_leaf_nodes': 2}  
5.578040812787403 {'max_features': 'auto', 'max_leaf_nodes': 6}  
6.392891119037453 {'max_features': 'auto', 'max_leaf_nodes': 10}  
6.497809150481358 {'max_features': 'auto', 'max_leaf_nodes': 14}  
6.181455805701661 {'max_features': 'auto', 'max_leaf_nodes': 16}  
6.619182266339363 {'max_features': 'auto', 'max_leaf_nodes': 20}  
5.1387572516143285 {'max_features': 'log2', 'max_leaf_nodes': 2}  
6.597922520142185 {'max_features': 'log2', 'max_leaf_nodes': 6}
```

```

7.126487423549713 {'max_features': 'log2', 'max_leaf_nodes': 10}
7.279222216155243 {'max_features': 'log2', 'max_leaf_nodes': 14}
7.905956094199112 {'max_features': 'log2', 'max_leaf_nodes': 16}
6.578257676718738 {'max_features': 'log2', 'max_leaf_nodes': 20}

```

In [67]: *#perform grid search again to find optimal max_depth*

```

gbr_param_grid = [
    {'max_leaf_nodes': [2],
     'max_features': ['log2'],
     'max_depth': [2,6,8,10,12]}]

gbr = GradientBoostingRegressor(n_estimators=100, learning_rate=1.0)

gbr_grid_search = GridSearchCV(gbr, gbr_param_grid, cv=5,
                                scoring='neg_mean_squared_error',
                                return_train_score=True)

gbr_grid_search.fit(boston_X, boston_y)

```

```

Out[67]: GridSearchCV(cv=5, error_score='raise',
                      estimator=GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=
                      learning_rate=1.0, loss='ls', max_depth=3, max_features=None,
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      n_estimators=100, presort='auto', random_state=None,
                      subsample=1.0, verbose=0, warm_start=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid=[{'max_leaf_nodes': [2], 'max_features': ['log2'], 'max_depth': [2,
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_squared_error', verbose=0)

```

```

In [68]: print('\n-----')
          print('The parameters chosen by Grid Search for Random Forest Regressor:\n ')
          gbr_grid_search.best_params_

```

```

-----
The parameters chosen by Grid Search for Random Forest Regressor:

```

```

Out[68]: {'max_depth': 8, 'max_features': 'log2', 'max_leaf_nodes': 2}

```

In [69]: gbr_cvres = gbr_grid_search.cv_results_

```

for mean_score, params in zip(gbr_cvres["mean_test_score"], gbr_cvres["params"]):
    print(np.sqrt(-mean_score), params)

```

```

6.272890234256745 {'max_depth': 2, 'max_features': 'log2', 'max_leaf_nodes': 2}
5.690240798235128 {'max_depth': 6, 'max_features': 'log2', 'max_leaf_nodes': 2}
5.439973992590656 {'max_depth': 8, 'max_features': 'log2', 'max_leaf_nodes': 2}
5.596202448488737 {'max_depth': 10, 'max_features': 'log2', 'max_leaf_nodes': 2}
5.8480637169228125 {'max_depth': 12, 'max_features': 'log2', 'max_leaf_nodes': 2}

```

0.8 Evaluate all models

0.8.1 Use KFold cross-validation

```

In [70]: # specify the set of classifiers being evaluated
         from sklearn.metrics import mean_squared_error

         names = ["Linear_Regression", "Ridge_Regression",
                  "Lasso", "Elastic_Net",
                  "Random_Forest", "Gradient_Boosting"]
         classifiers = [LinearRegression(),
                        Ridge(),
                        Lasso(),
                        ElasticNet(alpha=0.1, l1_ratio=0.25),
                        RandomForestRegressor(n_estimators=100, max_leaf_nodes=20,
                                             max_depth=8, n_jobs=-1, bootstrap=True,
                                             random_state=RANDOM_SEED),
                        GradientBoostingRegressor(n_estimators=100, max_depth=10,
                                                  learning_rate=1.0, max_features='log2',
                                                  max_leaf_nodes=2)]

In [71]: # -----
         # specify the k-fold cross-validation design
         from sklearn.model_selection import KFold

         # ten-fold cross-validation employed here
         N_FOLDS = 10

         # set up numpy array for storing results
         cv_results = np.zeros((N_FOLDS, len(names)))

In [72]: # Instantiate K-Folds cross-validator
         kf = KFold(n_splits = N_FOLDS, shuffle=False, random_state = RANDOM_SEED)

         # check the splitting process by looking at fold observation counts
         index_for_fold = 0 # fold count initialized
         for train_index, test_index in kf.split(model_data):
             print('\nFold index:', index_for_fold,
                   '-----')

         # note that 0:model_data.shape[1]-1 slices for explanatory variables
         # and model_data.shape[1]-1 is the index for the response variable
         X_train2 = model_data[train_index, 0:model_data.shape[1]-1]

```

```

X_test2 = model_data[test_index, 0:model_data.shape[1]-1]
y_train2 = model_data[train_index, model_data.shape[1]-1]
y_test2 = model_data[test_index, model_data.shape[1]-1]
print('\nShape of input data for this fold:',
      '\nData Set: (Observations, Variables)')
print('X_train:', X_train2.shape)
print('X_test:', X_test2.shape)
print('y_train:', y_train2.shape)
print('y_test:', y_test2.shape)

index_for_method = 0 # initialize
for name, clf in zip(names, classifiers):
    print('\nClassifier evaluation for:', name)
    print(' Scikit Learn method:', clf)
    clf.fit(X_train2, y_train2) # fit on the train set for this fold
    # evaluate on the test set for this fold
    y_test_predict = clf.predict(X_test2)
    fold_method_result = mean_squared_error(y_test2, y_test_predict)
    fold_rmse = np.sqrt(fold_method_result)
    print('RMSE:', fold_rmse)
    cv_results[index_for_fold, index_for_method] = fold_rmse
    index_for_method += 1

index_for_fold += 1

```

Fold index: 0 -----

Shape of input data for this fold:
Data Set: (Observations, Variables)
X_train: (455, 12)
X_test: (51, 12)
y_train: (455,)
y_test: (51,)

Classifier evaluation for: Linear_Regression

Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.6153307366878622

Classifier evaluation for: Ridge_Regression

Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None, normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.6160421022579456

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.9313027028555269

Classifier evaluation for: Elastic_Net

Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.6339112262359792

Classifier evaluation for: Random_Forest

Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8, max_features='auto', max_leaf_nodes=20, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1, oob_score=False, random_state=1, verbose=0, warm_start=False)

RMSE: 0.5787975457925774

Classifier evaluation for: Gradient_Boosting

Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=1.0, loss='ls', max_depth=10, max_features='log2', max_leaf_nodes=2, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)

RMSE: 0.5124946266766438

Fold index: 1 -----

Shape of input data for this fold:

Data Set: (Observations, Variables)

X_train: (455, 12)

X_test: (51, 12)

y_train: (455,)

y_test: (51,)

Classifier evaluation for: Linear_Regression

Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

RMSE: 0.3747267699016482

Classifier evaluation for: Ridge_Regression

Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None, normalize=False, random_state=None, solver='auto', tol=0.001)

RMSE: 0.3724873699458746

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,

```
normalize=False, positive=False, precompute=False, random_state=None,
selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.7831242969799976
```

Classifier evaluation for: Elastic_Net

```
Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25,
max_iter=1000, normalize=False, positive=False, precompute=False,
random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.3323423970650091
```

Classifier evaluation for: Random_Forest

```
Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8,
max_features='auto', max_leaf_nodes=20,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
oob_score=False, random_state=1, verbose=0, warm_start=False)
RMSE: 0.31801072119037505
```

Classifier evaluation for: Gradient_Boosting

```
Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
learning_rate=1.0, loss='ls', max_depth=10,
max_features='log2', max_leaf_nodes=2,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100,
presort='auto', random_state=None, subsample=1.0, verbose=0,
warm_start=False)
RMSE: 0.35198734019035294
```

Fold index: 2 -----

Shape of input data for this fold:

Data Set: (Observations, Variables)

X_train: (455, 12)

X_test: (51, 12)

y_train: (455,)

y_test: (51,)

Classifier evaluation for: Linear_Regression

```
Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.7372119896892206
```

Classifier evaluation for: Ridge_Regression

```
Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.7361441517989507
```

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 1.1137912300446424

Classifier evaluation for: Elastic_Net

Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.718563165173973

Classifier evaluation for: Random_Forest

Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8, max_features='auto', max_leaf_nodes=20, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1, oob_score=False, random_state=1, verbose=0, warm_start=False)

RMSE: 0.5150723852566453

Classifier evaluation for: Gradient_Boosting

Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=1.0, loss='ls', max_depth=10, max_features='log2', max_leaf_nodes=2, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)

RMSE: 0.7865314510515283

Fold index: 3 -----

Shape of input data for this fold:

Data Set: (Observations, Variables)

X_train: (455, 12)

X_test: (51, 12)

y_train: (455,)

y_test: (51,)

Classifier evaluation for: Linear_Regression

Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

RMSE: 0.6535643910371687

Classifier evaluation for: Ridge_Regression

Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None, normalize=False, random_state=None, solver='auto', tol=0.001)

RMSE: 0.6524778781987847

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.9651492613503816

Classifier evaluation for: Elastic_Net

Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.5656235706545696

Classifier evaluation for: Random_Forest

Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8, max_features='auto', max_leaf_nodes=20, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1, oob_score=False, random_state=1, verbose=0, warm_start=False)

RMSE: 0.39089290645481534

Classifier evaluation for: Gradient_Boosting

Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=1.0, loss='ls', max_depth=10, max_features='log2', max_leaf_nodes=2, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)

RMSE: 0.7925112929211118

Fold index: 4 -----

Shape of input data for this fold:

Data Set: (Observations, Variables)

X_train: (455, 12)

X_test: (51, 12)

y_train: (455,)

y_test: (51,)

Classifier evaluation for: Linear_Regression

Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

RMSE: 0.6534559101742444

Classifier evaluation for: Ridge_Regression


```
Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
    normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.6528224169003206
```

Classifier evaluation for: Lasso

```
Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
    normalize=False, positive=False, precompute=False, random_state=None,
    selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.9343170336152345
```

Classifier evaluation for: Elastic_Net

```
Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25,
    max_iter=1000, normalize=False, positive=False, precompute=False,
    random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.5949392103369477
```

Classifier evaluation for: Random_Forest

```
Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8,
    max_features='auto', max_leaf_nodes=20,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
    oob_score=False, random_state=1, verbose=0, warm_start=False)
RMSE: 0.6394416159467092
```

Classifier evaluation for: Gradient_Boosting

```
Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
    learning_rate=1.0, loss='ls', max_depth=10,
    max_features='log2', max_leaf_nodes=2,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    presort='auto', random_state=None, subsample=1.0, verbose=0,
    warm_start=False)
RMSE: 0.6402966920976941
```

Fold index: 5 -----

```
Shape of input data for this fold:
Data Set: (Observations, Variables)
X_train: (455, 12)
X_test: (51, 12)
y_train: (455,)
y_test: (51,)
```

Classifier evaluation for: Linear_Regression

```
Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.4204278869595973
```

Classifier evaluation for: Ridge_Regression

Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None, normalize=False, random_state=None, solver='auto', tol=0.001)

RMSE: 0.4194772785083912

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.9521726271573456

Classifier evaluation for: Elastic_Net

Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

RMSE: 0.3973994633664465

Classifier evaluation for: Random_Forest

Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8, max_features='auto', max_leaf_nodes=20, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1, oob_score=False, random_state=1, verbose=0, warm_start=False)

RMSE: 0.3632158461863273

Classifier evaluation for: Gradient_Boosting

Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=1.0, loss='ls', max_depth=10, max_features='log2', max_leaf_nodes=2, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)

RMSE: 0.3977960825437818

Fold index: 6 -----

Shape of input data for this fold:

Data Set: (Observations, Variables)

X_train: (456, 12)

X_test: (50, 12)

y_train: (456,)

y_test: (50,)

Classifier evaluation for: Linear_Regression

Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.4044402074544354

Classifier evaluation for: Ridge_Regression

Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None, normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.4044339063399666

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.7600693575647395

Classifier evaluation for: Elastic_Net

Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.44973435865315786

Classifier evaluation for: Random_Forest

Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8, max_features='auto', max_leaf_nodes=20, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1, oob_score=False, random_state=1, verbose=0, warm_start=False)
RMSE: 0.47327136560355426

Classifier evaluation for: Gradient_Boosting

Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=1.0, loss='ls', max_depth=10, max_features='log2', max_leaf_nodes=2, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)
RMSE: 0.473629434783943

Fold index: 7 -----

Shape of input data for this fold:
Data Set: (Observations, Variables)
X_train: (456, 12)
X_test: (50, 12)
y_train: (456,)
y_test: (50,)

Classifier evaluation for: Linear_Regression

Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.8044974066807946

Classifier evaluation for: Ridge_Regression

Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None, normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.8054744909557435

Classifier evaluation for: Lasso

Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 1.5513050279342468

Classifier evaluation for: Elastic_Net

Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25, max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.8365288425953381

Classifier evaluation for: Random_Forest

Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8, max_features='auto', max_leaf_nodes=20, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1, oob_score=False, random_state=1, verbose=0, warm_start=False)
RMSE: 0.6809513919693294

Classifier evaluation for: Gradient_Boosting

Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=1.0, loss='ls', max_depth=10, max_features='log2', max_leaf_nodes=2, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)
RMSE: 0.8952207358858596

Fold index: 8 -----

Shape of input data for this fold:

Data Set: (Observations, Variables)

X_train: (456, 12)

X_test: (50, 12)

```
y_train: (456,)
y_test: (50,)
```

Classifier evaluation for: Linear_Regression

```
Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.6890695563753627
```

Classifier evaluation for: Ridge_Regression

```
Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
    normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.6898685677369694
```

Classifier evaluation for: Lasso

```
Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
    normalize=False, positive=False, precompute=False, random_state=None,
    selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 1.4387124078783402
```

Classifier evaluation for: Elastic_Net

```
Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25,
    max_iter=1000, normalize=False, positive=False, precompute=False,
    random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.7277584919693205
```

Classifier evaluation for: Random_Forest

```
Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8,
    max_features='auto', max_leaf_nodes=20,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
    oob_score=False, random_state=1, verbose=0, warm_start=False)
RMSE: 0.6920356402367644
```

Classifier evaluation for: Gradient_Boosting

```
Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
    learning_rate=1.0, loss='ls', max_depth=10,
    max_features='log2', max_leaf_nodes=2,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    presort='auto', random_state=None, subsample=1.0, verbose=0,
    warm_start=False)
RMSE: 0.781336644228036
```

Fold index: 9 -----

Shape of input data for this fold:
Data Set: (Observations, Variables)

```
X_train: (456, 12)
X_test: (50, 12)
y_train: (456,)
y_test: (50,)
```

Classifier evaluation for: Linear_Regression

```
Scikit Learn method: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
RMSE: 0.47386940618795154
```

Classifier evaluation for: Ridge_Regression

```
Scikit Learn method: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
    normalize=False, random_state=None, solver='auto', tol=0.001)
RMSE: 0.47172152353606756
```

Classifier evaluation for: Lasso

```
Scikit Learn method: Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
    normalize=False, positive=False, precompute=False, random_state=None,
    selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.7617856516578413
```

Classifier evaluation for: Elastic_Net

```
Scikit Learn method: ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.25,
    max_iter=1000, normalize=False, positive=False, precompute=False,
    random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
RMSE: 0.4662848152155337
```

Classifier evaluation for: Random_Forest

```
Scikit Learn method: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=8,
    max_features='auto', max_leaf_nodes=20,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
    oob_score=False, random_state=1, verbose=0, warm_start=False)
RMSE: 0.43554980135162846
```

Classifier evaluation for: Gradient_Boosting

```
Scikit Learn method: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
    learning_rate=1.0, loss='ls', max_depth=10,
    max_features='log2', max_leaf_nodes=2,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    presort='auto', random_state=None, subsample=1.0, verbose=0,
    warm_start=False)
RMSE: 0.5373245439259907
```

```
In [73]: # convert array into a data frame
```

```

# and then assign column names from names list
cv_results_df = pd.DataFrame(cv_results)
cv_results_df.columns = names

print('\n-----')
print('Average results from ', N_FOLDS, '-fold cross-validation\n',
      '\nMethod          RMSE', sep = '')
print(cv_results_df.mean())

```

Average results from 10-fold cross-validation

Method	RMSE
Linear_Regression	0.582659
Ridge_Regression	0.582095
Lasso	1.019173
Elastic_Net	0.572309
Random_Forest	0.508724
Gradient_Boosting	0.616913

dtype: float64

```

In [74]: print('\n-----')
print('RMSE for each fold of cross validation:\n')
print(cv_results_df)

```

RMSE for each fold of cross validation:

	Linear_Regression	Ridge_Regression	Lasso	Elastic_Net	Random_Forest	\
0	0.615331	0.616042	0.931303	0.633911	0.578798	
1	0.374727	0.372487	0.783124	0.332342	0.318011	
2	0.737212	0.736144	1.113791	0.718563	0.515072	
3	0.653564	0.652478	0.965149	0.565624	0.390893	
4	0.653456	0.652822	0.934317	0.594939	0.639442	
5	0.420428	0.419477	0.952173	0.397399	0.363216	
6	0.404440	0.404434	0.760069	0.449734	0.473271	
7	0.804497	0.805474	1.551305	0.836529	0.680951	
8	0.689070	0.689869	1.438712	0.727758	0.692036	
9	0.473869	0.471722	0.761786	0.466285	0.435550	
	Gradient_Boosting					
0	0.512495					
1	0.351987					
2	0.786531					
3	0.792511					

4	0.640297
5	0.397796
6	0.473629
7	0.895221
8	0.781337
9	0.537325

```
In [75]: #From Geron (2017) Hands-On Machine Learning with Scikit-Learn and TensorFlow, Chapter 8
#Yet another great quality of Random Forests is that they make it
#easy to measure the relative importance of each feature.
#Scikit-Learn measures a features importance by looking at how
#much the tree nodes that use that feature reduce impurity on average
 #(across all trees in the forest). More precisely, it is a weighted average,
#where each nodes weight is equal to the number of training samples
#that are associated with it.

#Scikit-Learn computes this score automatically for each feature
#after training, then it scales the results so that the sum of all importances is equal to 1

#Generate the final model with optimal parameter settings
#and view the relative importance of each feature

rfr_final = RandomForestRegressor(n_estimators=100, max_leaf_nodes=20,
                                max_depth=8, n_jobs=-1, bootstrap=True, random_state=RANDOM_SEED)
rfr_final.fit(boston_X, boston_y)
print('Relative weights of each predictor:')
for name, score in zip(boston_X.columns, rfr_final.feature_importances_):
    print(name, score)
```

Relative weights of each predictor:

```
crim 0.027186900064009843
zn 0.00013358541161817637
indus 0.0033123533343774754
chas 0.0004578441552995733
nox 0.021057327392709692
rooms 0.4448076728101155
age 0.004430352909061474
dis 0.06641855285178194
rad 0.0015801670327649446
tax 0.010703377686566025
ptratio 0.015617060678483929
lstat 0.4042948056732115
```